

Concurrency

Paolo Burgio

paolo.burgio@unimore.it



UNIMORE
UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA

High Performance
Real Time **Lab**

FRATELLI

QUANDO LA MAMMA ESCE DI CASA



QUANDO TORNA A CASA





Why concurrency?

Functional

- › Many users may be connected to the same system at the same time
- › Each user can have its own processes that execute concurrently with the processes of the other users
- › Perform many operations concurrently

Performance

- › Take advantage of blocking time
- › While some thread waits for a blocking condition, another thread performs another operation
- › On a multi-core machine, independent activities can be carried out on different cores at the same time



Cooperative vs. Competitive

Competitive concurrency

- › Different activities compete for the resources
- › One activity does not know anything about the other
- › The OS must manage the resources so to
 - Avoid conflicts
 - Be fair

Cooperative concurrency

- › Many activities cooperate to perform an operation
- › Every activity knows about the others
- › They must synchronize on particular events



Competitive

Competing activities need to be “protected” from each other

- › Separate memory spaces, as with different processes

The **allocation** of the resource and the synchronization must be **centralized**

- › Competitive activities request for services to a central manager (the OS or some dedicated process) which allocates the resources in a fair way

Client/Server model

- › Communication is usually done through **messages**

More suitable to the **process** model of execution



Client/server model

A server manages the resource **exclusively**


- › For example, the printer

If a process needs to access the resource, it sends a **request to the server**

- › For example, printing a file, or asking for the status
- › The server can send back the responses
- › The server can also be on a remote system

Two basic primitives:

- › **send** and **receive**





Cooperative model

Cooperative activities **know** about each other

- › They do not need memory protection (less overhead)

They need to access the same data structures

- › Allocation of the resource is **de-centralized**
- › **Shared memory** model

More suitable to the **thread** model of execution



Competition vs. cooperation

Competition is best resolved by using the **message passing** model

- › However it can be implemented using a shared memory paradigm too

Cooperation is best implemented by using the **shared memory** paradigm

- › However, it can be realized by using pure message passing mechanisms

General purpose OS needs to support both models

- › Protection for competing activities
- › Client/server models → message passing primitives
- › Shared memory for reducing the overhead

Some special OS supports only one of the two

- › RTOS supports only shared memory



Models of concurrency



Message passing

Message passing systems are based on the basic concept of message

Two basic operations

- send(destination, message);
 - send can be synchronous or asynchronous (*fire-and-forget*)
- receive(source, &message);
 - receive can be symmetric or asymmetric



Producer/Consumer with MP

- The producer executes `send(consumer, data)`
- the consumer executes `receive(producer, data)`
- no need for a special communication structure (already contained in the send/receive semantic)





Resources and message passing

There are no shared resources in the message passing model
–all the resources are allocated statically, accessed in a dedicated way

Each resource is handled by a **manager process** that is the only one that has right to access to a resource

The consistency of a data structure is guaranteed by the manager process
–there is no more competition, only cooperation!!!



Synchronous communication

synchronous send/receive

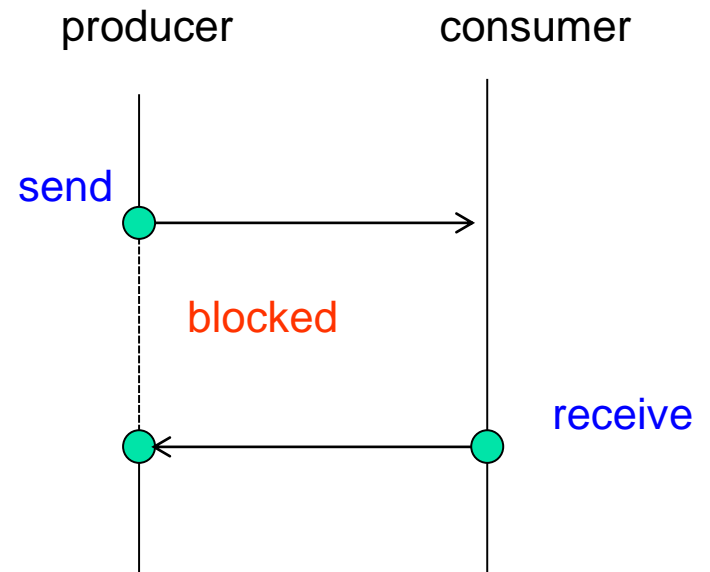

–no buffers!

producer:

```
s_send(consumer, d);
```

consumer:

```
s_receive(producer, &d);
```





Async send/ sync receive

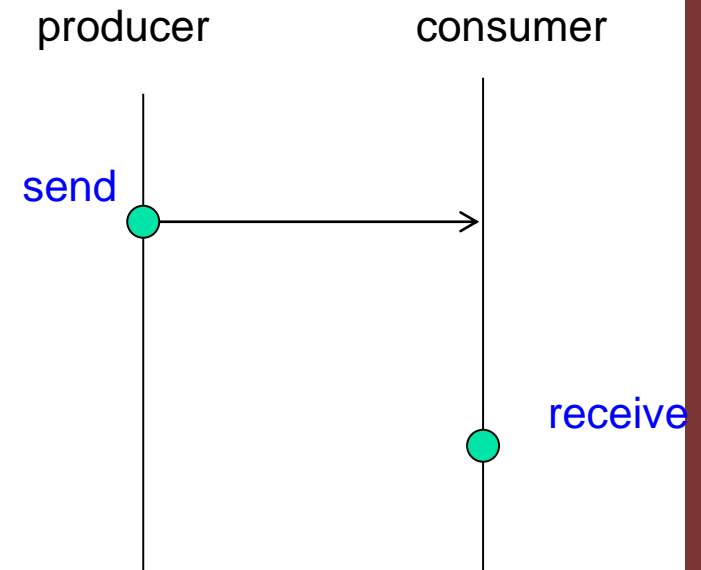

asynchronous send / synchronous receive
-there is probably a send buffer somewhere

producer:

```
a_send(consumer, d);
```

consumer:

```
s_receive(producer, &d);
```





(A)symmetric receive

Symmetric receive

- receive(source, &data);
- the programmer wants a message from a given producer

Asymmetric receive

- source = receive(&data);
- often, we do not know who is the sender

- imagine a web server;
- the programmer cannot know in advance the address of the browser that will request the service
- many browsers can ask for the same service




Remote procedure call

Increase expressiveness

From low-level (MP) to more «programmer friendly» (RPC) mechanism

- In a client-server system, a client wants to request an action to a server –that is typically done using a **remote procedure call (RPC)**





Message passing systems

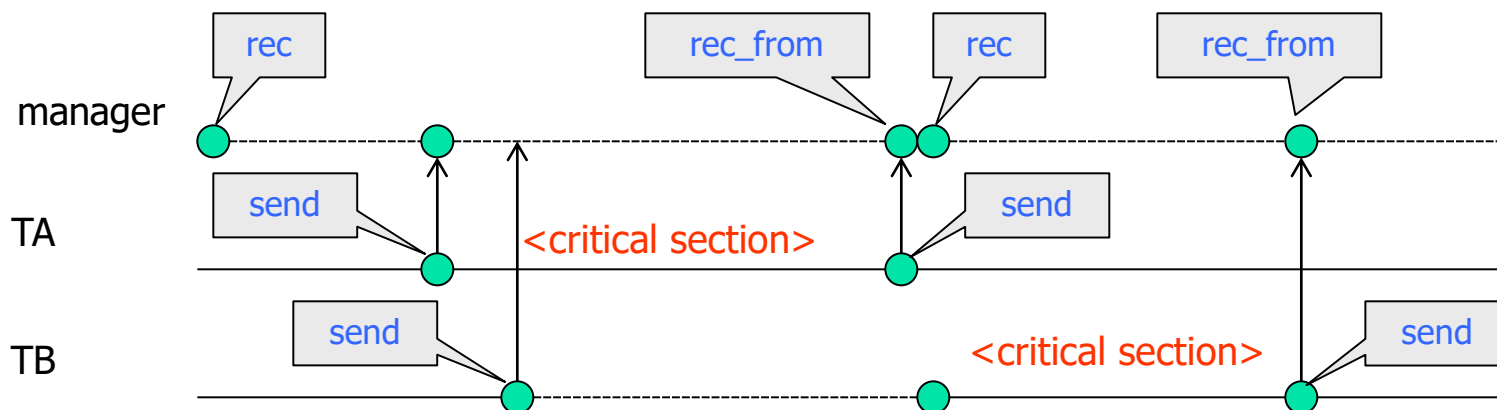
- In message passing
 - each resource needs one threads manager (often called **daemon thread**)
 - the threads manager is responsible for giving access to the resource
- Example: let's try to implement mutual exclusion with message passing primitives
 - one thread will ensure mutual exclusion
 - every thread that wants to access the resource must
 - send a message to the manager thread
 - access the critical section
 - send a message to signal the leaving of the critical section



Sync send / sync receive

```
void * manager(void *)  
{  
    thread_t source;  
    int d;  
    while (true) {  
        source = s_receive(&d);  
        s_receive_from(source, &d);  
    }  
}
```

```
void * thread(void *)  
{  
    int d;  
    while (true) {  
        s_send(manager, d);  
        <critical section>  
        s_send(manager, d);  
    }  
}
```

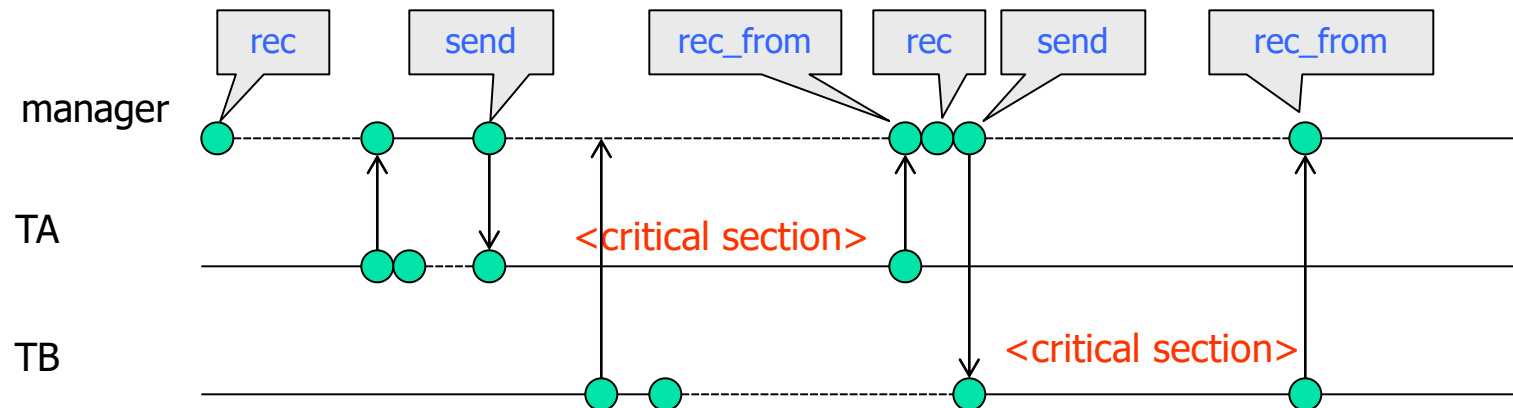




With async send and sync receive

```
void * manager(void *)
{
    thread_t source;
    int d;
    while (true) {
        source = s_receive(&d);
        a_send(source,d);
        s_receive_from(source,&d);
    }
}
```

```
void * thread(void *)
{
    int d;
    while (true) {
        a_send(manager, d);
        s_receive_from(manager, &d);
        <critical section>
        a_send(manager, d);
    }
}
```





Problem

- Implement readers/writers with message passing
- Hints:
 - define a manager thread
 - the service type (read/write) can be passed as data
 - use asynchronous send and synchronous receive
 - use symmetric and asymmetric receive




Shared memory model



Shared memory

- › The first one being supported in old OS
- › The simplest one and the closest to the machine
- › All threads can access the same memory locations






Hardware analogy

An abstract model that presents a good analogy is the following

- › Many HW CPU, each one running one activity (thread)
- › One shared memory





Resource allocation

- Allocation of resource can be
 - **Static**: once the resource is granted, it is never revoked
 - **Dynamic**: resource can be granted and revoked dynamically
 - Manager
- Access to a resource can be
 - **Dedicated**: only one activity at a time may request access to the resource
 - **Shared**: many activities may access the resource at the same time
 - Mutual exclusion

	Dedicated	Shared
Static	Compile Time	Manager
Dynamic	Manager	Manager



Mutual exclusion problem

- We do not know in advance the relative speed of the processes
 - Hence, we do not know the order of execution of the hardware instructions
- Example:
 - Incrementing a variable x is NOT an atomic operation



Atomicity

- A hardware instruction is atomic if it cannot be “interleaved” with other instructions
 - Atomic operations are always sequentialized
 - Atomic operations cannot be interrupted
- They are safe operations
 - For example, transferring one word from memory to register or viceversa
 - Non atomic operations can be interrupted
- They are not “safe” operations
- Non elementary operations are not atomic



Non-atomic operations

- Consider a “simple” operation like:

>

```
x = x+1;
```

- In assembler:

>

```
LD R0, x  
INC R0  
ST x, R0
```

- A simple operation like incrementing a memory variable, may be composed by three machine instructions

>



Example 1

- Bad interleaving:
shared memory

```
int x ;
```

```
void *threadA(void *)  
{  
    ...;  
    x = x + 1;  
    ...;  
}
```

```
void *threadB(void *)  
{  
    ...;  
    x = x + 1;  
    ...;  
}
```

	...		
LD R0, x	TA	x = 0	
LD R0, x	TB	x = 0	
INC R0	TB	x = 0	
ST x, R0	TB	x = 1	
INC R0	TA	x = 1	
ST x, R0	TA	x = 1	
	...		



Example 2

- Bad interleaving
Shared object (sw resource)

```
struct A_t {  
    int a;  
    int b;  
} A;  
  
void A_init(A_t *x) { x->a=1;      x->b=1; }  
void A_inc(A_t *x) { x->a++;      x->b++; }  
void A_mul(A_t *x) { x->b*=2;    x->a*=2; }
```

```
void *threadA(void *)  
{  
    ...  
    A_inc(&A);  
    ...  
}
```

```
void *threadB(void *)  
{  
    ...  
    A_mul(&A);  
    ...  
}
```

x->a++;	TA	a = 2
x->b*=2;	TB	b = 2
x->b++;	TA	b = 3
x->a*=2;	TB	a = 4

consistency:
after each
operation,
a == b

resource in a
non-consistent
state!



Consistency

- For each resource, we can state some **consistency property**
 - A consistency property C_i is a **boolean expression** on the values of the **internal variables**
 - A consistency property must hold **before** and **after** each operation
 - It does **not** hold **during an operation**
 - If the operations are properly sequentialized, the consistency properties must hold
- **Formal verification**
 - Let R be a resource, and let $C(R)$ be a set of consistency properties on the resource
 - $C(R) = \{ C_i \}$
 - › **Definition:** a concurrent program is **correct** if, for every possible interleaving of the operations on the resource, the consistency properties hold after each operation





Example 3: circular array

```
struct CircularArray_t {
    int array[10];
    int head, tail, num;
} queue;

void init_CA(struct CircularArray_t *a)
{ a->head=0; a->tail=0; a->num=0; }

int insert_CA(struct CircularArray_t *a,
              int elem)
{
    if (a->num == 10) return 0;
    a->array[a->head] = elem;
    a->head = (a->head + 1) % 10;
    a->num++;
    return 1;
}

int extract_CA(struct CircularArray_t *a,
               int *elem)
{
    if (a->num == 0) return 0;
    *elem = a->array[a->tail];
    a->tail = (a->tail + 1) % 10;
    a->num--;
    return 1;
}

/* suppose num++ e num- atomic */
```

Consistency properties

(suppose num++ and num-- atomic)

C_1 : if (num == 0 || num == 10)
head == tail;

C_2 : if (0 < num < 10)
num == (head - tail) % 10


C_3 : num == NI - NE

C_4 : (insert x)
pre: if (num < 10)
post: num == num + 1 &&
array[(head-1)%10] = x;

C_5 : (extract &x)
pre: if (num > 0)
post: num == num - 1 &&
x = array[(tail-1)%10];



Example 3: circular array - insert



Initial state:

head = 0; tail = 0; num = 0;

insert_CA (&queue, 5) ;

head = 1; tail = 0; num = 1;

C_2, C_3, C_4
holds


C_2 : if ($0 < \text{num} < 10$)
 $\text{num} == (\text{head} - \text{tail}) \% 10$

C_3 : $\text{num} == NI - NE$

C_4 : insert_CA(&queue, x)
pre: if ($\text{num} < 10$)
post: $\text{num} == \text{num} + 1 \ \&\&$
 $\text{array}[(\text{head}-1)\%10] = x;$



Example 3: circular array – insert (2)



Initial state:

head = 0; tail = 0; num = 0;

insert_CA (&queue, 5) ;

head = 1; tail = 0; num = 1;

insert_CA (&queue, 3) ;

head = 2; tail = 0; num = 2;

C_2, C_3, C_4
hold


C_2 : if ($0 < \text{num} < 10$)
 $\text{num} == (\text{head} - \text{tail}) \% 10$

C_3 : $\text{num} == \text{NI} - \text{NE}$

C_4 : insert_CA(&queue, x)
pre: if ($\text{num} < 10$)
post: $\text{num} == \text{num} + 1 \ \&\&$
 $\text{array}[(\text{head}-1)\%10] = x;$



Example 3: circular array – insert ⁽³⁾



Initial state:

$head = 9; tail = 5; num = 4;$

$insert_CA (&queue, 6);$

$head = 0; tail = 5; num = 5$

C_2, C_3, C_4
hold


$C_2:$ $if (0 < num < 10)$
 $num == (head - tail) \% 10$

$C_3:$ $num == NI - NE$

$C_4:$ $insert_CA (&queue, x)$
pre: $if (num < 10)$
post: $num == num + 1 \ \&\&$
 $array[(head-1)\%10] = x;$



Example 3: circular array – extract



Initial state:

$\text{head} = 0; \text{tail} = 5; \text{num} = 5;$

$\text{extract_CA} (\&\text{queue}, \&\text{elem}) ;$

$\text{head} = 0; \text{tail} = 6; \text{num} = 4$

$C_2:$ $\text{if } (0 < \text{num} < 10)$
 $\text{num} == (\text{head} - \text{tail}) \% 10$

$C_3:$ $\text{num} == \text{NI} - \text{NE}$

$C_5:$ $\text{extract_CA} (\&\text{queue}, \&\text{x})$
 pre: $\text{if } (\text{num} > 0)$
 post: $\text{num} == \text{num} - 1 \ \&\&$
 $\text{x} = \text{array}[\text{tail}];$


C_2, C_3, C_5
hold



Example 3: the problem

- If the insert operation is performed by two processes, some consistency property may be violated!

```
struct CircularArray_t queue;
```




```
void *threadA(void *)  
{  
    ...  
    insert_CA( &queue, 5);  
    ...  
}
```

```
void *threadB(void *)  
{  
    ...  
    insert_CA( &queue, 2);  
    ...  
}
```



Example 3: interference



head (**)

head (*)

C₄ is violated!

5 != array[head - 1]

```
if (a->num == 10) return 0;
a->array[a->head] = 5;
a->head = (a->head + 1) % 10; (**)
a->num ++;
return 1;
```

Initial state:

head = 7; tail = 3; num = 4;

insert_CA (&queue, 5) ; (TA)

insert_CA (&queue, 2) ; (TB)

```
if (a->num == 10) return 0;      (TA)
a->array[a->head] = 5;            (TA)
if (a->num == 10) return 0;      (TB)
a->array[a->head] = 2;            (TB)
a->head = (a->head + 1) % 10; (TB) (*)
a->num ++;                      (TB)
return 1;                      (TB)
a->head = (a->head + 1) % 10; (TA) (**)
a->num ++;                      (TA)
return 1;                      (TA)
```

```
if (a->num == 10) return 0;
a->array[a->head] = 2;
a->head = (a->head + 1) % 10; (*)
a->num ++;
return 1;
```

Final State:

head = 9; tail = 3; num = 6;



Example 3: correctness

- The previous program is **not correct**
 - It exist a possible interleaving of two insert operations that leaves the resource in a inconsistent state
- Proving the non-correctness is easy
 - it suffices to find a counter example
- Proving the correctness is not easy
 - it is necessary to prove the correctness for every possible interleaving of every operation



Example 3: problem

- › What if an insert and an extract are interleaved?
 - Nothing bad can happen!!
 - Proof
- › if $0 < \text{num} < 10$, insert_CA() and extract_CA() are independent
- › if $\text{num} == 0$
 - if extract_CA begins before insert_CA, it immediately returns 0, so nothing bad can happen
 - if insert_CA begins before, extract_CA will either return false, or it will find a new element to extract, depending on how it interleaves w.r.t. $\text{num}++$ of insert_CA (last operation of insert_CA)
- › Dual thing when $\text{num} == 10$



Example 3: CircularArray properties

- **a)** if more than one thread executes insert_CA()
 - inconsistency!!
- **b)** if we have only two threads
 - one threads calls insert_CA() and the other thread calls extract_CA()
 - no inconsistency!
- The order of the operations is important!
 - a wrong order can make the object inconsistency even under the assumption b)
 - the case is when num is incremented but the data has not yet been inserted
 - in any case, the final result depends on the timings of the dfferent requests (e.g, an insertion with the buffer full)



Example 3: questions

- Problem:

- In the previous example, we supposed that `num++` and `num--` are atomic operations
- What happens if they are not atomic?

- Question:

- Assuming that operation `--` and `++` are not atomic, can we make the `circularArray` safe under the assumption b) ?
- Hint: try to substitute variable `num` with two boolean variables: `bool empty` and `bool full`;



Critical sections

- Definitions

- The **shared object** where the conflict may happen is a “**resource**”
- The **parts of the code** where the problem may happen are called “**critical sections**”
- A critical section is a sequence of operations that cannot be interleaved with other operations on the same resource
- Two critical sections on the same resource must be properly sequentialized
- We say that two critical sections on the same resource must execute in **MUTUAL EXCLUSION**
- There are two ways to obtain mutual exclusion
 - Disabling the preemption** (valid only for single-core systems)
 - Implementing the critical section as an **atomic operation**, using **semaphores** and **mutexes**




Critical sections: disabling preemption

- Single core systems

- In some scheduler, it is possible to **disable preemption** for a limited interval of time

- Problems:

- If a **high priority critical thread needs to execute**, it cannot make preemption and it is delayed
 - Even if the high priority task does not access the resource!





Critical sections: atomic operations

- There exist some general mechanisms to implement mutual exclusion only between the processes that uses a resource:

–semaphores

–mutexes

- Define a **flag s** for each resource
- Use **lock(s)/unlock(s)** around the critical section

```
int s;  
...  
lock(s);  
<critical section>  
unlock(s);  
...
```



Synchronisation

- Mutual exclusion is not the only problem
- We need a **way of synchronise two or more threads**
- Example: **producer/consumer**
- Suppose we have two threads,
 - One produces some integers and sends them to another thread (PRODUCER)
 - Another one takes the integer and elaborates it (CONSUMER)





Producer/consumer

- The two threads have different speeds
 - For example the producer is much faster than the consumer
 - We need to store the integers in a queue, so that no data is lost
 - Let's use the `CircularArray_t` structure



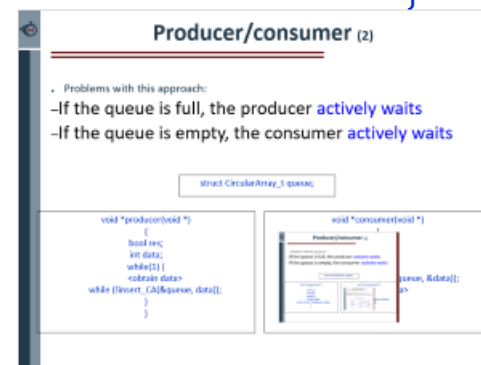
Producer/consumer (2)

- Problems with this approach:
 - If the queue is full, the producer **actively waits**
 - If the queue is empty, the consumer **actively waits**

```
struct CircularArray_t queue;
```

```
void *producer(void *)  
{  
    bool res;  
    int data;  
    while(1) {  
        <obtain data>  
        while (!insert_CA(&queue, data));  
    }  
}
```

```
void *consumer(void *)
```



```
queue, &data));  
>
```



A more general approach

- We need to provide a general mechanism for synchronisation and mutual exclusion
- Requirements
 - Provide mutual exclusion between critical sections
 - Avoid two insertions operation to interleave
 - Synchronise two threads on one condition
 - For example, block the producer when the queue is full



General mechanism: semaphores

- Dijkstra proposed the semaphore mechanism
- A semaphore is an abstract entity that consists of
 - A counter
 - A blocking queue
 - Operation wait
 - Operation signal
- The operations on a semaphore are considered atomic



Semaphores

- Semaphores are basic mechanisms for providing synchronization
 - It has been shown that every kind of synchronization and mutual exclusion can be implemented by using semaphores
 - We will analyze possible implementation of the semaphore mechanism later

```
typedef struct {  
    <blocked queue> blocked;  
    int counter;  
} sem_t;  
  
void sem_init      (sem_t *s, int n);  
  
void sem_wait      (sem_t *s);  
void sem_post      (sem_t *s);
```

Note:
the real prototype
of sem_init is
slightly different!



Wait and signal

- A **wait** operation has the following behavior
 - If counter == 0, the requiring thread is blocked
 - It is removed from the ready queue
 - It is inserted in the blocked queue
 - If counter > 0, then counter--;
- A **post** operation has the following behavior
 - If counter == 0 and there is some blocked thread, unblock it
 - The thread is removed from the blocked queue
 - It is inserted in the ready queue
 - Otherwise, increment counter



Semaphores

```
void sem_init (sem_t *s, int n)
{
    s->count=n;
    ...
}

void sem_wait(sem_t *s)
{
    if (counter == 0)
        <block the thread>
    else
        counter--;
}

void sem_post(sem_t *s)
{
    if (<there are blocked threads>)
        <unblock a thread>
    else
        counter++;
}
```



Signal semantics

- What happens when a thread blocks on a semaphore?
 - In general, it is inserted in a BLOCKED queue
- Extraction from the blocking queue can follow different semantics:
 - Strong semaphore
 - The threads are removed in well-specified order
 - For example, FIFO order, priority based ordering, ...
 - Signal and suspend
 - After the new thread has been unblocked, a thread switch happens
 - Signal and continue
 - After the new thread has been unblocked, the thread that executed the signal continues to execute
- Concurrent programs should not rely too much on the semaphore semantic



Mutual exclusion with semaphores

- How to use a semaphore for critical sections
 - Define a semaphore **initialized to 1**
 - Before entering the critical section, perform a **wait**
 - After leaving the critical section, perform a **post**

```
sem_t s;  
...  
sem_init(&s, 1);
```


```
void *threadA(void *arg)  
{  
    ...  
    sem_wait(&s);  
    <critical section>  
    sem_post(&s);  
    ...  
}
```

```
void *threadB(void *arg)  
{  
    ...  
    sem_wait(&s);  
    <critical section>  
    sem_post(&s);  
    ...  
}
```



Mutual exclusion with semaphores (2)

semaphore

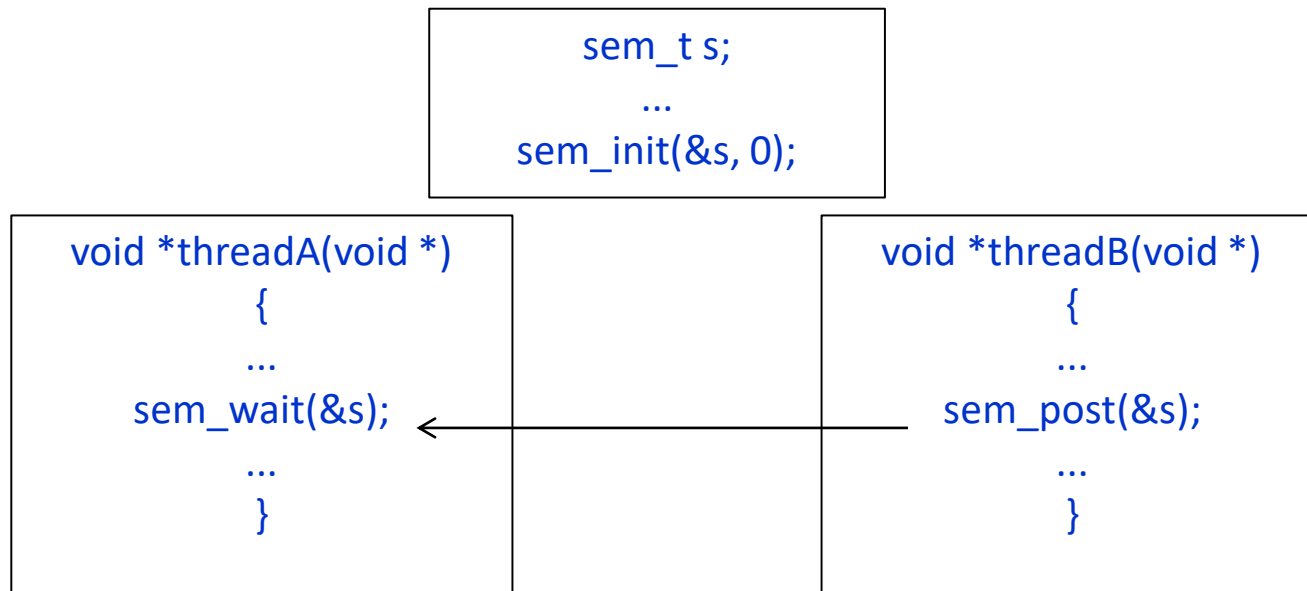


<code>sem_wait();</code>	(TA)
<code><critical section (1)></code>	(TA)
<code>sem_wait()</code>	(TB)
<code><critical section (2)></code>	(TA)
<code>sem_post()</code>	(TA)
<code><critical section></code>	(TB)
<code>sem_post()</code>	(TB)



Synchronization

- How to use a semaphore for synchronization
 - Define a semaphore **initialized to 0**
 - At the synchronization point, follower performs a **wait**
 - At the synchronization point, producer performs a **post**
 - In the example, threadA blocks until threadB wakes it up



- How can both A and B synchronize at the same point?



Producer/consumer

- Consider a producer/consumer system
 - One producer executes `insert_CA()`
 - We want the producer to be blocked when the queue is full
 - The producer will be unblocked when there is some space again
 - One consumer executes `extract_CA()`
 - We want the consumer to be blocked when the queue is empty
 - The consumer will be unblocked when there is some space again
 - First attempt: one producer and one consumer only



Producer/consumer (2)

```
struct CircularArray_t {
    int array[10];
    int head, tail;
    sem_t empty, full;
}

void init_CA(struct CircularArray_t *c)
{ c->head=0; c->tail=0;
  sem_init(&c->empty, 0); sem_init(&c->full, 10); }

void insert_CA(struct CircularArray_t *c, int elem) {
    sem_wait(&c->full);
    c->array[c->head] = elem;
    c->head = (c->head + 1) % 10;
    sem_post(&c->empty);
}

void extract_CA(struct CircularArray_t *c, int &elem) {
    sem_wait(&c->empty);
    elem = c->array[c->tail];
    c->tail = (c->tail + 1) % 10;
    sem_post(&c->full);
}
```

Note: there is no member called *num* as we had in Exa



Producer/consumer: properties

- Notice that

- The value of the counter of **empty** is the **number of elements** in the queue
 - It is the number of times we can call extract without blocking
- The value of the counter of **full** is the complement of the elements in the queue
 - It is the number of times we can call insert without blocking

- Exercise

- Prove that the implementation is correct
 - insert_CA() never overwrites elements
 - extract_CA() always gets an element of the queue



Producers/consumers

- Now let's combine mutual exclusion and synchronization
 - Consider a system in which there are
 - Many producers
 - Many consumers
 - We want to implement synchronization
 - We want to protect the data structure



Producers/consumers: does it work?

```
struct CircularArray_t {
    int array[10];
    int head, tail;
    sem_t full, empty;
    sem_t mutex;
}

void init_CA(struct CircularArray_t *c)
{
    c->head=0; c->tail=0;
    sem_init(&c->empty, 0); sem_init(&c->full, 10); sem_init(&c->mutex, 1);
}
```

```
void insert_CA(struct CircularArray_t *c,
               int elem)
{
    sem_wait(&c->mutex);
    sem_wait(&c->full);
    c->array[c->head]=elem;
    c->head = (c->head+1)%10;
    sem_post(&c->empty);
    sem_post(&c->mutex);
}
```

```
void extract_CA(struct CircularArray_t *c,
                int *elem)
{
    sem_wait(&c->mutex);
    sem_wait(&c->empty);
    elem = c->array[c->tail];
    c->tail = (c->tail+1)%10;
    sem_post(&c->full);
    sem_post(&c->mutex);
}
```



Producers/consumers: correct solution

```
struct CircularArray_t {
    int array[10];
    int head, tail;
    sem_t full, empty;
    sem_t mutex;
}

void init_CA(struct CircularArray_t *c)
{
    c->head=0; c->tail=0;
    sem_init(&c->empty, 0); sem_init(&c->full, 10); sem_init(&c->mutex, 1);
}
```

```
void insert_CA(struct CircularArray_t *c,
               int elem)
{
    sem_wait(&c->full);
    sem_wait(&c->mutex);
    c->array[c->head]=elem;
    c->head = (c->head+1)%10;
    sem_post(&c->mutex);
    sem_post(&c->empty);
}
```

```
void extract_CA(struct CircularArray_t *c,
                int *elem)
{
    sem_wait(&c->empty);
    sem_wait(&c->mutex);
    elem = c->array[c->tail];
    c->tail = (c->tail+1)%10;
    sem_post(&c->mutex);
    sem_post(&c->full);
}
```



Producers/consumers: deadlock situation

- Deadlock situation

- A thread executes `sem_wait(&c->mutex)` and then blocks on a synchronisation semaphore
- To be unblocked another thread must enter a critical section guarded by the same mutex semaphore!
- So, the first thread cannot be unblocked and free the mutex!
- The situation **cannot be solved**, and the two threads will never proceed

- As a rule, **never insert a blocking synchronization inside a critical section!!!**



Readers/writers

- One shared buffer
- Readers:
 - They read the content of the buffer
 - Many readers can read at the same time
- Writers
 - They write in the buffer
 - While one writer is writing no other reader or writer can access the buffer
- Use semaphores to implement the resource



Readers/writers: simple implementation

```
struct Buffer_t {  
    sem_t synch;  
    sem_t s_R;  
    int nr;  
}  
  
void init_B(struct Buffer_t *b)  
{ sem_init(&b->synch, 1);  
  sem_init(&b->s_R, 1);  
  b->nr=0; }
```

```
void read_B(struct Buffer_t *b) {  
    sem_wait(&b->s_R);  
    b->nr++;  
    if (b->nr==1) sem_wait(&b->synch);  
    sem_post(&b->s_R);  
  
    <read the buffer>  
  
    sem_wait(&b->s_R);  
    b->nr--;  
    if (b->nr==0) sem_post(&b->synch);  
    sem_post(&b->s_R);  
}
```

```
void write_B(struct Buffer_t *b) {  
    sem_wait(&b->synch);  
  
    <write the buffer>  
  
    sem_post(&b->synch);  
}
```



Readers/writers: more than one pending writer

```
struct Buffer_t {  
    sem_t synch, mutex;  
    sem_t s_R, s_W;  
    int nr, nw;  
};
```

```
void init_B(struct Buffer_t *b)  
{  
    sem_init(&b->synch, 1); sem_init(&b->mutex(1);  
    sem_init(&b->s_R, 1); sem_init(&b->s_W, 1);  
    b->nr=0; b->nw=0;  
}
```

```
void read_B(struct Buffer_t *b) {  
    sem_wait(&b->s_R);  
    b->nr++;  
    if (b->nr==1)  
        sem_wait(&b->synch);  
    sem_post(&b->s_R);  
    <read the buffer>  
    sem_wait(&b->s_R);  
    b->nr--;  
    if (b->nr==0)  
        sem_post(&b->synch);  
    sem_post(&b->s_R);  
}
```

```
void write_B(struct Buffer_t *b) {  
    sem_wait(&b->s_W);  
    b->nw++;  
    if (b->nw==1) sem_wait(&b->synch);  
    sem_post(&b->s_W);  
  
    sem_wait(&b->mutex);  
    <write the buffer>  
    sem_post(&b->mutex);  
  
    sem_wait(&b->s_W);  
    b->nw--;  
    if (b->nw==0) sem_post(&b->synch);  
    sem_post(&b->s_W);  
}
```



Readers/writers: starvation

- A reader will be blocked for a finite time
- The writer suffers starvation
- Suppose we have 2 readers (R1 and R2) and 1 writer W1
 - Suppose that R1 starts to read
 - While R1 is reading, W1 blocks because it wants to write
 - R2 starts to read
 - R1 finishes, but, since R2 is reading, W1 cannot be unblocked
 - Before R2 finishes to read, R1 starts to read again
 - When R2 finishes, W1 cannot be unblocked because R1 is reading
- A solution
 - Readers should not be counted whenever there is a writer waiting for them



Readers/writers: priority to writers!

```
struct Buffer_t {  
    sem_t synch, synch1;  
    sem_t s_R, s_W;  
    int nr, nw;  
};
```

```
void init_B(struct Buffer_t *b) {  
    sem_init(&b->synch, 1); sem_init(&b->synch1, 1);  
    sem_init(&b->s_R, 1); sem_init(&b->s_W, 1);  
    b->nr=0; b->nw=0;  
}
```

```
void read_B(struct Buffer_t *b) {
```

```
    sem_wait(&b->synch1);
```

```
    sem_wait(&b->s_R);
```

```
    b->nr++;
```

```
    if (b->nr==1) sem_wait(&b->synch);
```

```
    sem_post(&b->s_R);
```

```
    sem_post(&b->synch1);
```

```
    <read the buffer>
```

```
    sem_wait(&b->s_R);
```

```
    b->nr--;
```

```
    if (b->nr==0) sem_post(&b->synch);
```

```
    sem_post(&b->s_R);
```

```
}
```

```
void write_B(struct Buffer_t *b) {
```

```
    sem_wait(&b->s_W);
```

```
    b->nw++;
```

```
    if (b->nw==1) sem_wait(&b->synch1);
```

```
    sem_post(&b->s_W);
```

```
    sem_wait(&b->synch);
```

```
    <write the buffer>
```

```
    sem_post(&b->synch);
```

```
    sem_wait(&b->s_W);
```

```
    b->nw--;
```

```
    if (b->nw == 0) sem_post(&b->synch1);
```

```
    sem_post(&b->s_W);
```

```
}
```



Readers/writers: problem

- Now, there is starvation for readers
- The readers/writers problem can be solved in general?
 - No starvation for readers
 - No starvation for writers
- Solution
 - Maintain a FIFO ordering with requests
 - If at least one writer is blocked, every next reader blocks
 - If at least one reader is blocked, every next writer blocks
- We can do that using the **private semaphores** technique

References



Course website

- › http://hipert.unimore.it/people/paolob/pub/Industrial_Informatics/index.html

My contacts

- › paolo.burgio@unimore.it
- › <http://hipert.mat.unimore.it/people/paolob/>

Resources

- › Giorgio Buttazzo, "Hard Real-Time Computing Systems : Predictable Scheduling Algorithms and Applications". 3rd Edition. 2011. Springer
- › "Real-Time Embedded Systems" course by Prof. Bertogna @UNIMORE
- › A "small blog"
 - <http://www.google.com>