

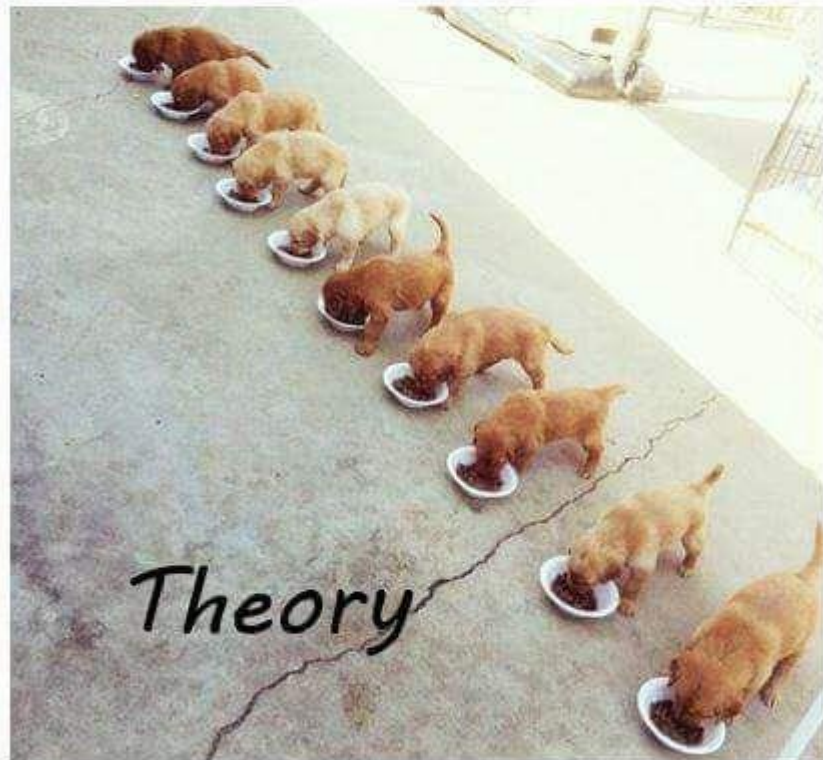
# Parallel programming

---

Paolo Burgio

paolo.burgio@unimore.it

# Multithreaded programming





# Definitions

---

## › Parallel computing

- Partition computation across different compute engines
- E.g., PThreads w/Shared mem, but also multi-process on the same machine!

## › Distributed computing

- Partition computation across different machines
- E.g., multiprocess (MPI, MQTT) w/message passing

Same principle, more general



# Why do we need parallel computing?

---

Increase performance of our machines

## › Scale-up

- Solve a "bigger" problem in the same time

## › Scale-out

- Solve the same problem in less time



# Yes but..

---

Why (highly) parallel machines...

...and not faster single-core machines?



# The answer #1 - Money

---





# The answer #2 – the "hot" one

---

## Moore's law

- › "The number of transistors that we can pack in a given die area doubles every 18 months"

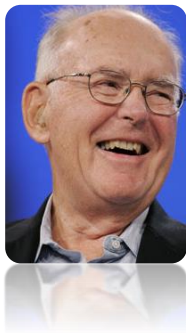
## Dennard's scaling

- › "performance per watt of computing is growing exponentially at roughly the same rate"



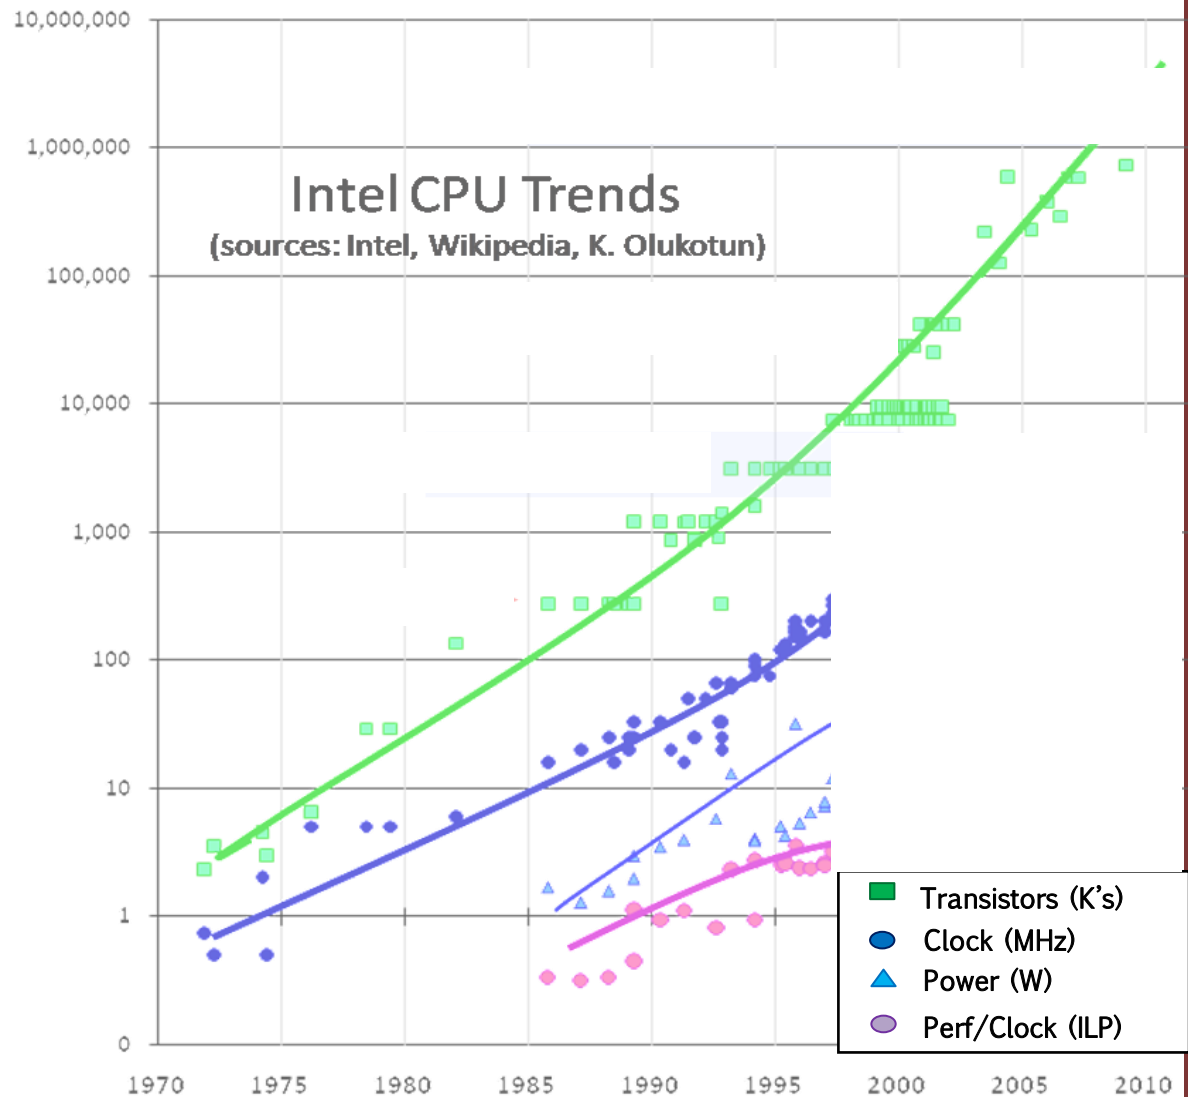
# The answer #2 – the "hot" one

## › SoC design paradigm



## › Gordon Moore

- His law is still valid, but...







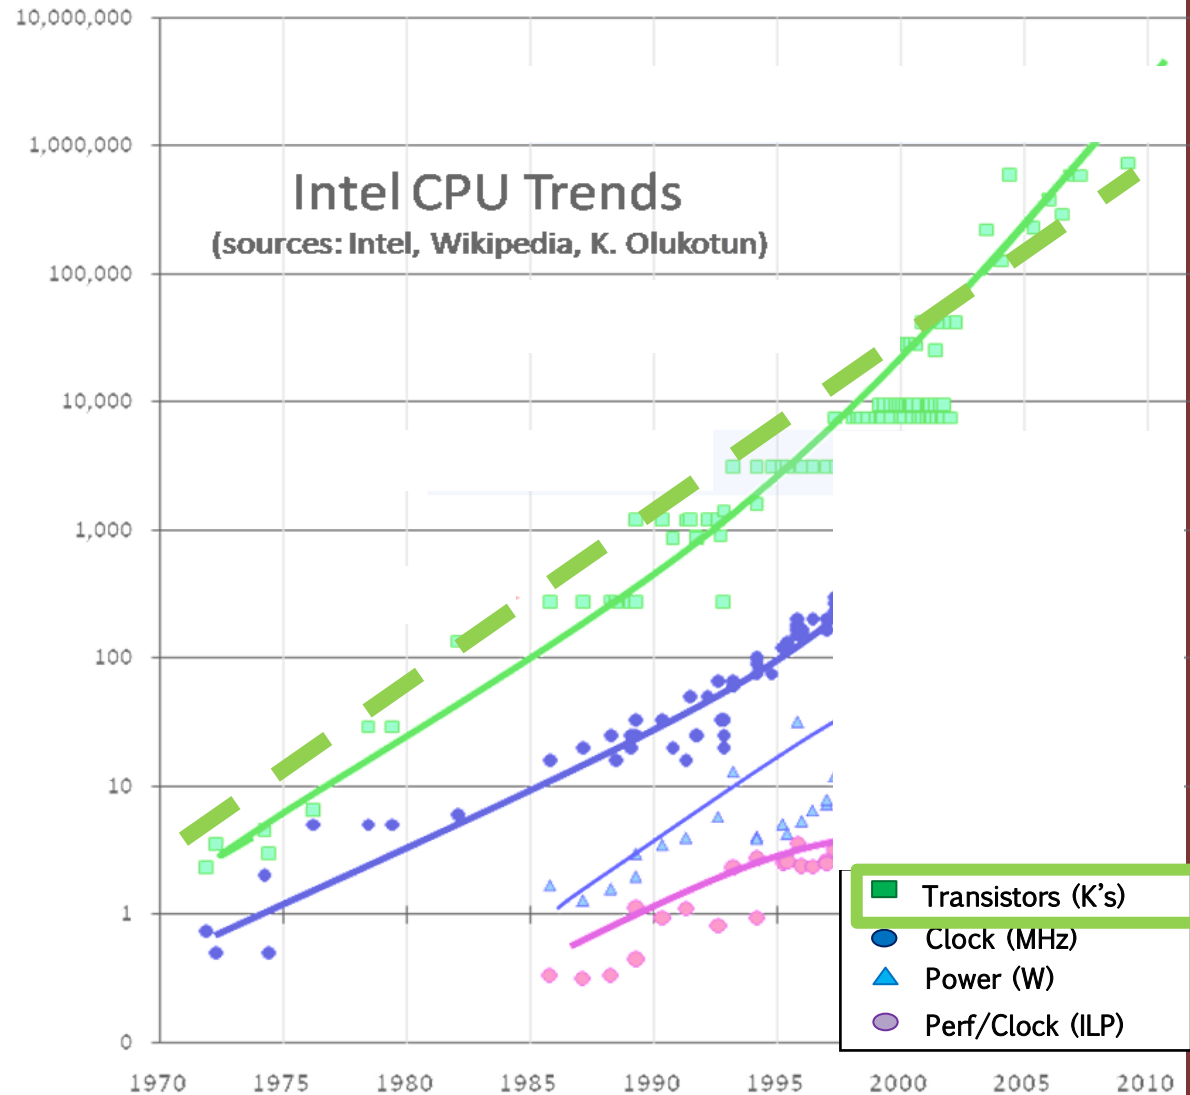
# The answer #2 – the "hot" one

## › SoC design paradigm



## › Gordon Moore

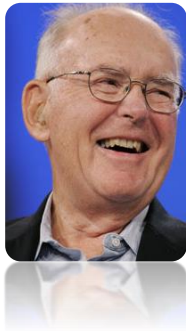
- His law is still valid, but...





# The answer #2 – the "hot" one

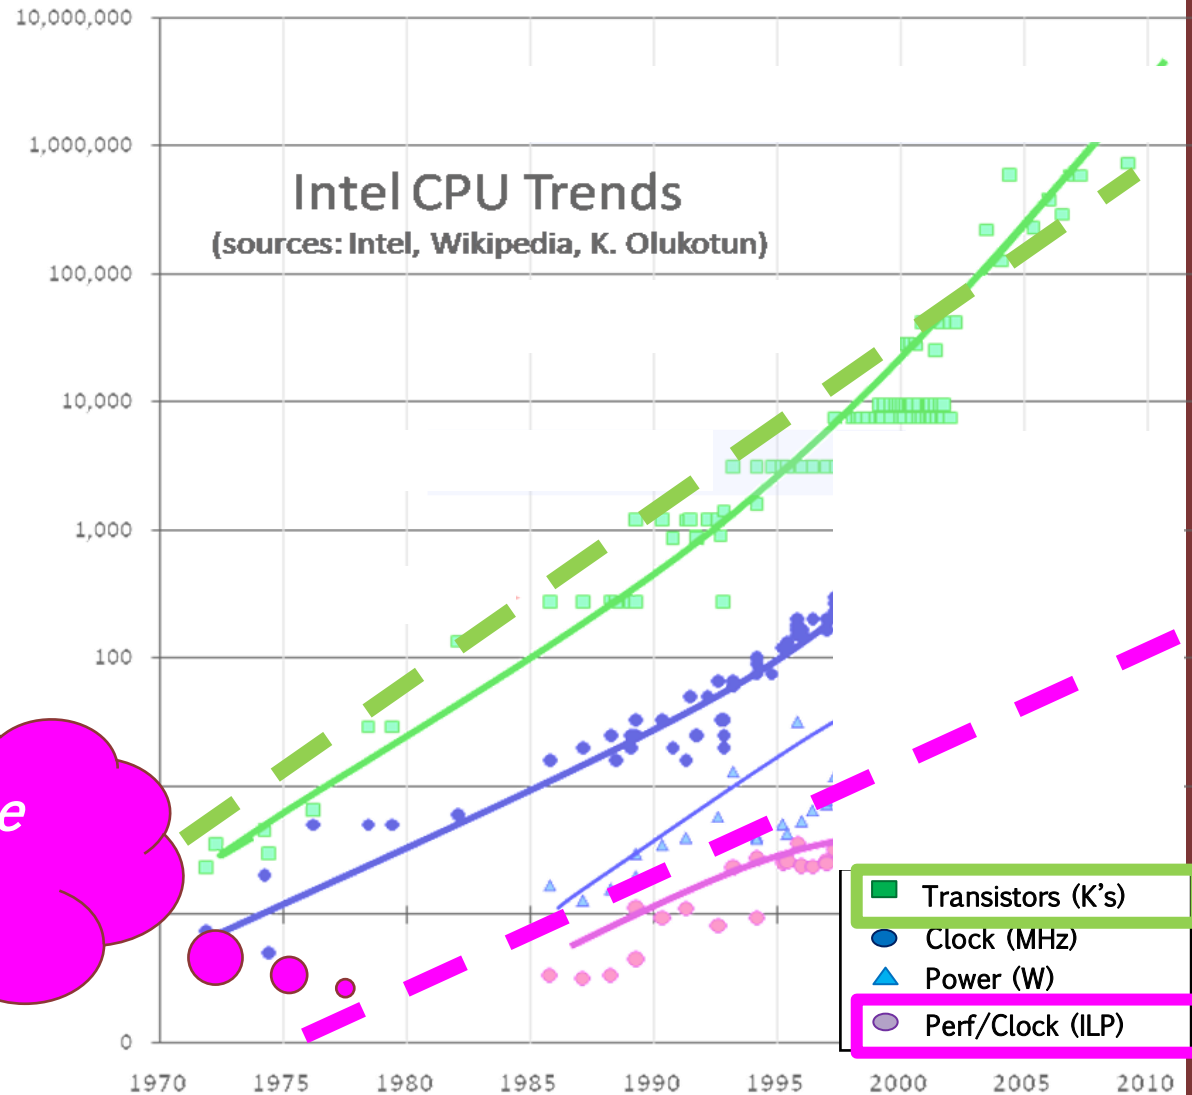
## › SoC design paradigm



## › Gordon Moore

- His law is still valid, but...

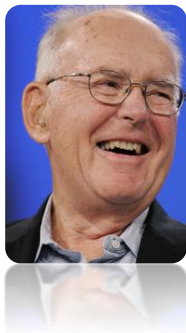
*Performance  
→  
frequency*





# The answer #2 – the "hot" one

## › SoC design paradigm



## › Gordon Moore

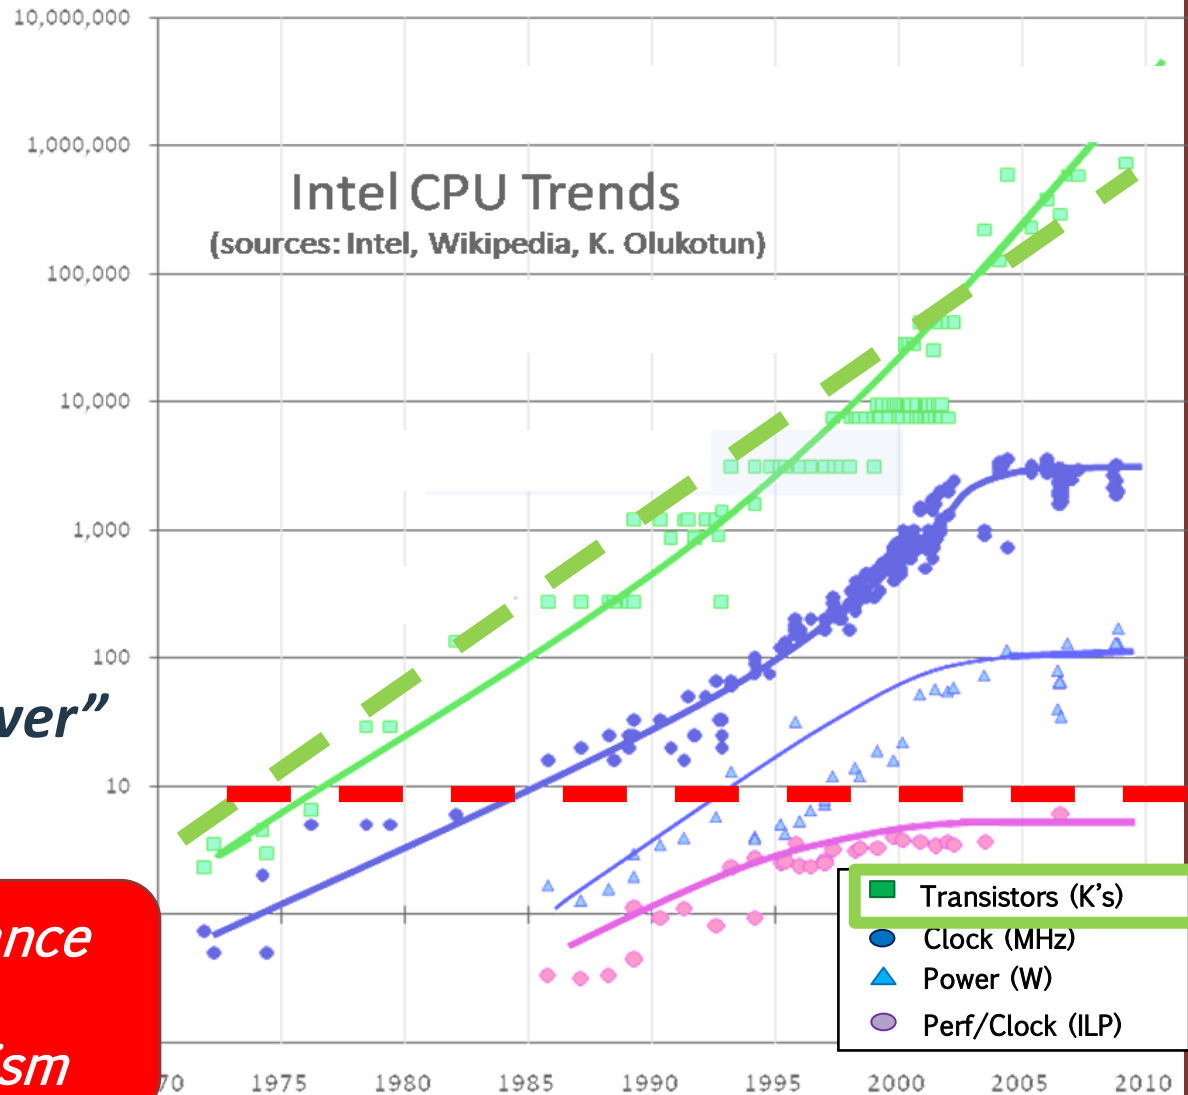
- His law is still valid, but...

## › “The free lunch is over”

- Herb Sutter, 2005

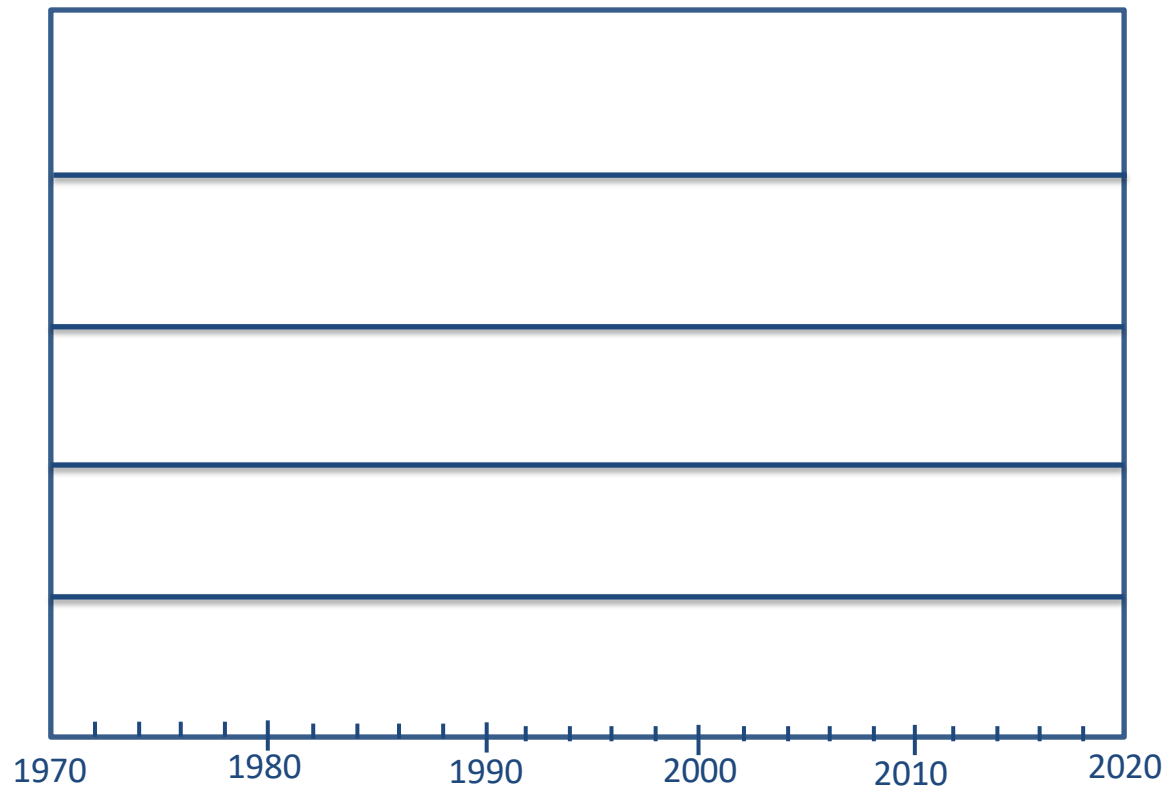


*Performance  
→  
parallelism*





# In other words...

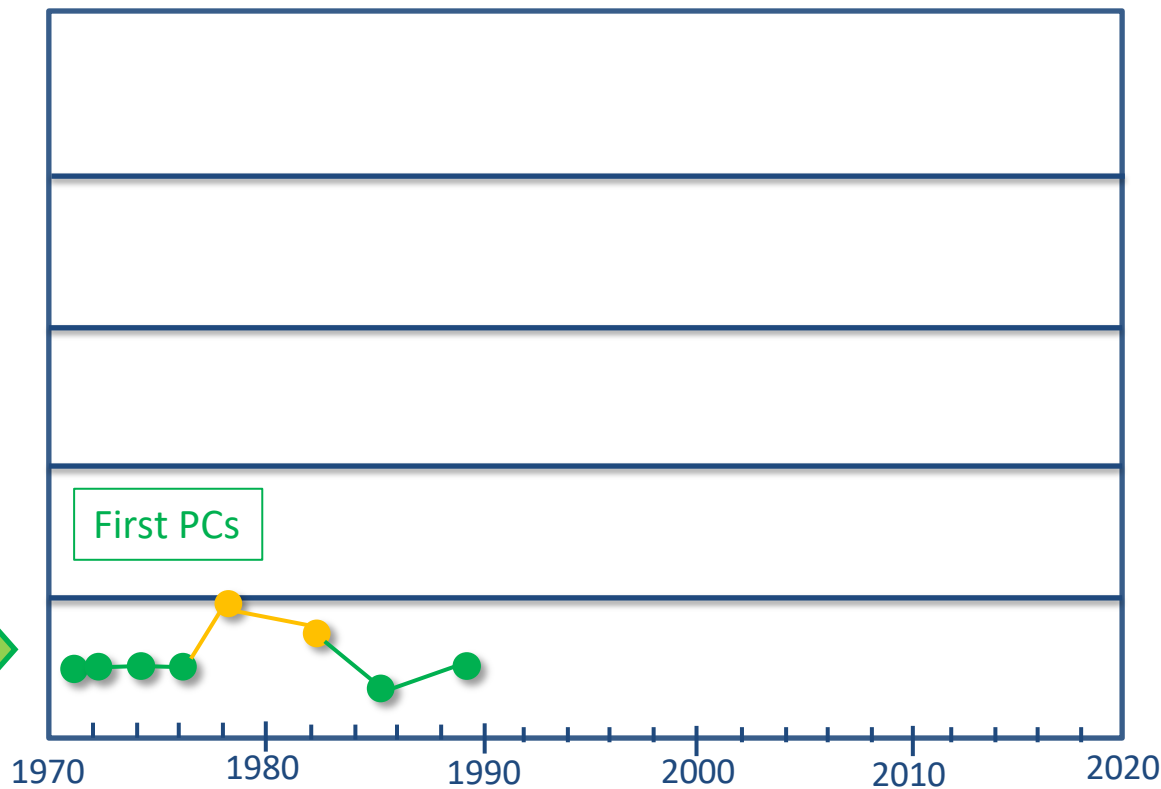




# In other words...



Summer  
temperature

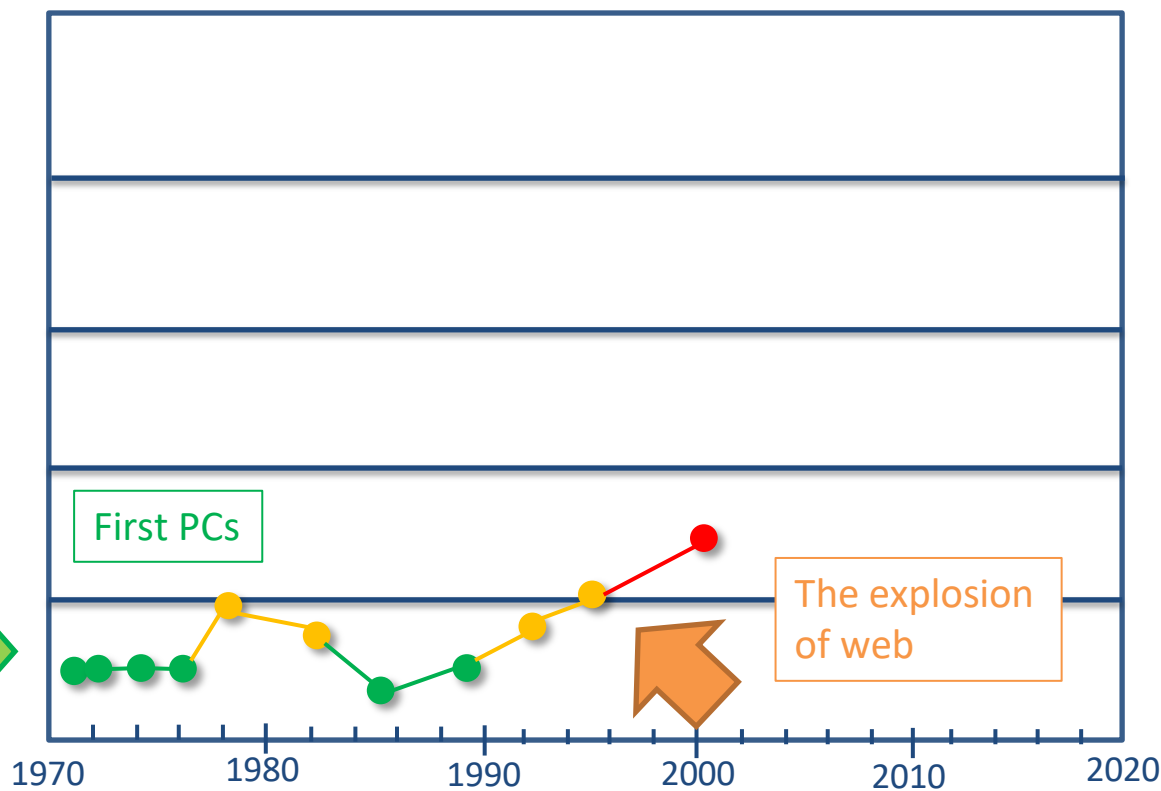




# In other words...



Summer  
temperature



Hot plate

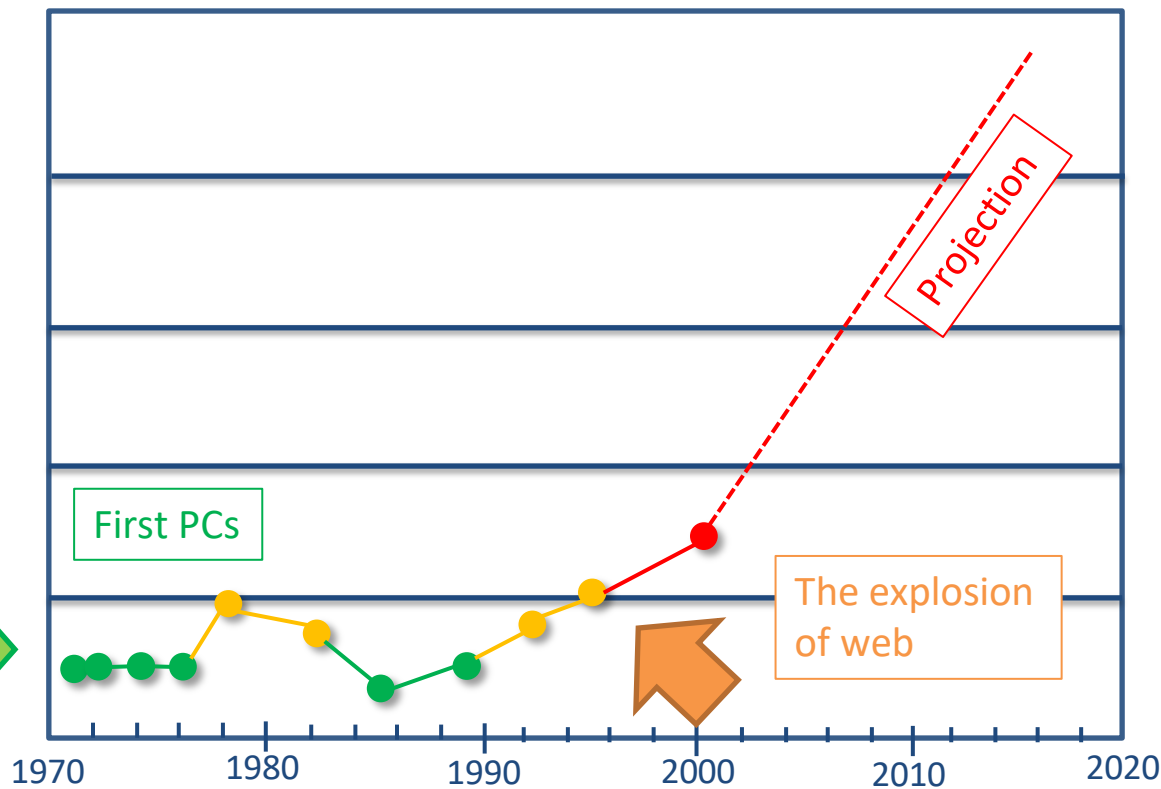




# In other words...



Summer  
temperature

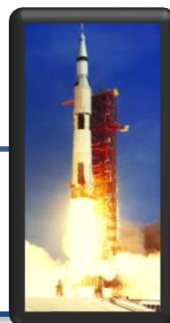


Hot plate

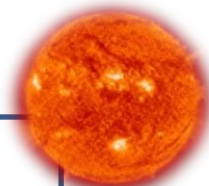




# In other words...



Surface of the sun



Rocket nozzle

Nuclear Reactor



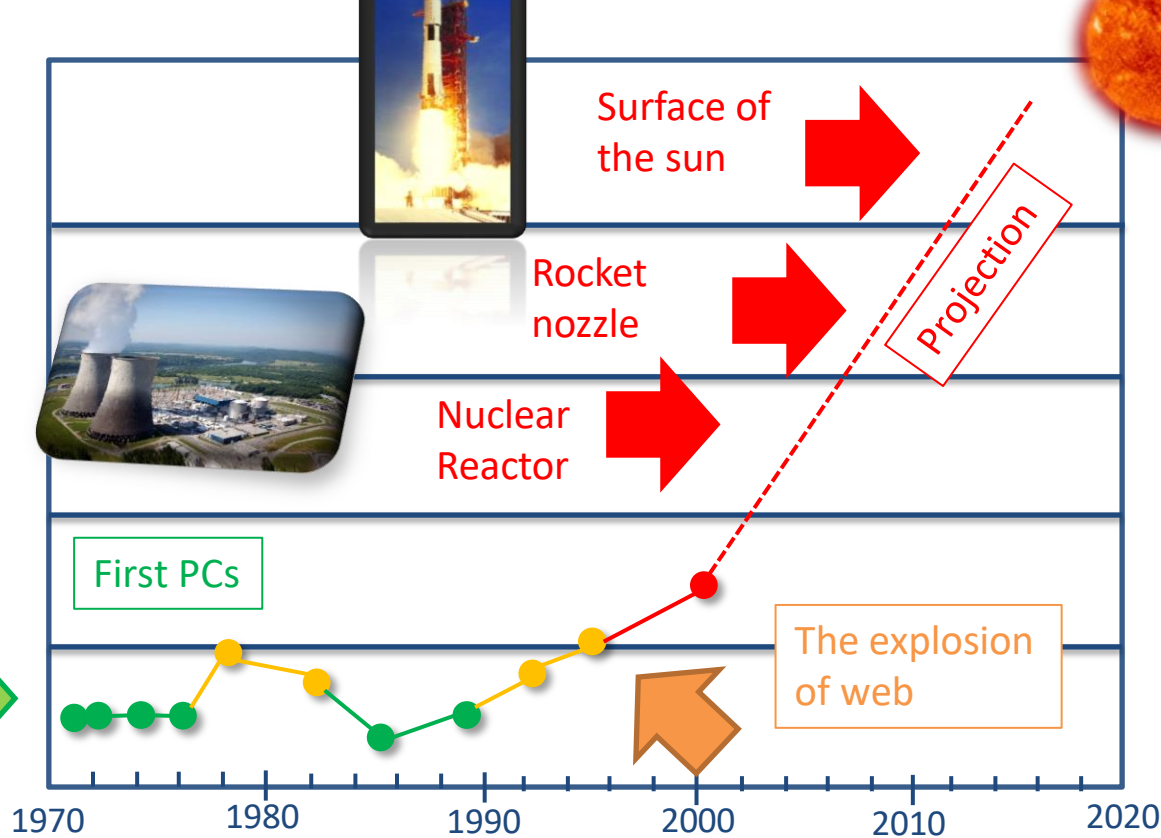
First PCs

The explosion of web

Hot plate



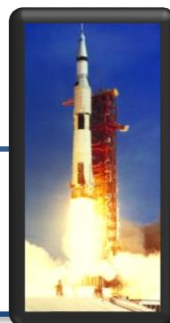
Summer temperature



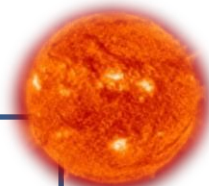




# In other words...



Surface of the sun



Rocket nozzle

Nuclear Reactor

Modern computers

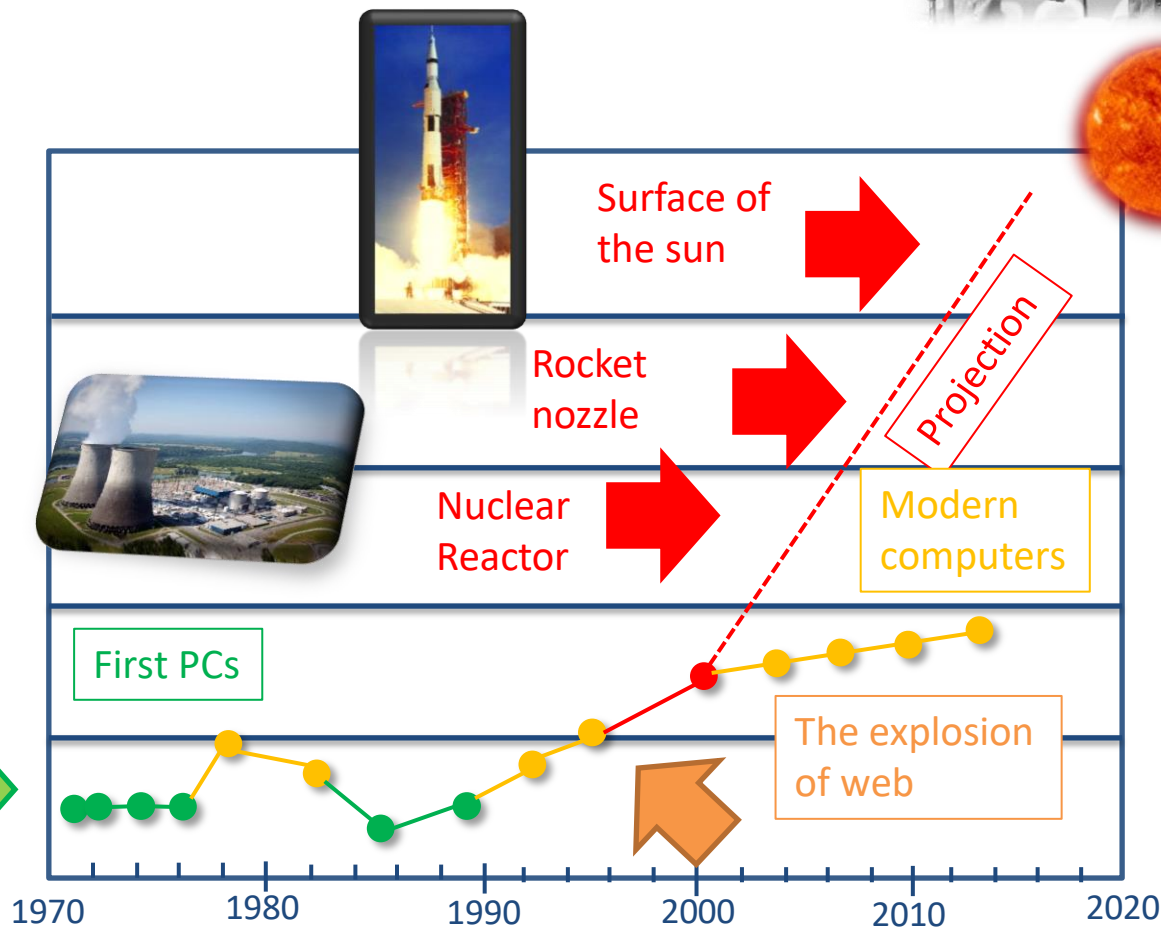
First PCs

The explosion of web

Hot plate



Summer temperature





# Instead of going faster..

---

- › ..(go faster but through) parallelism!

## Problem #1

- › New computer architectures
- › At least, three architectural templates

## Problem #2

- › Need to efficiently program them
- › HPC already has this problem!

## The problem

- › Programmers **must** know a bit of the architecture!
- › To make parallelization effective
- › "Let's run this on a GPU. It certainly goes faster" (cit.)

# The Big problem

- › Effectively programming in parallel is difficult

*Brian Kernighan (1942-)*

- *Researcher, theory of informatics*
- *Co-authored UNIX and AWK*
- *Wrote "The C Programming Language" book*

“Everyone knows that debugging is twice as hard as writing a program in the first place.

So if you're as clever as you can be when you write it, how will you ever debug it?”





# Amdahl's law

---

- › A sequential program that takes 100 sec to exec
- › Only 95% can run in parallel (it's a lot)
- › And.. you are an extremely good programmer, and you have a machine with 1billion cores, so that part takes 0 sec
- › So,

$$T_{par} = 100_{sec} - 95_{sec} = 5_{sec}$$

$$Speedup = \frac{100_{sec}}{5_{sec}} = 20x$$

...20x, on one billion cores!!!

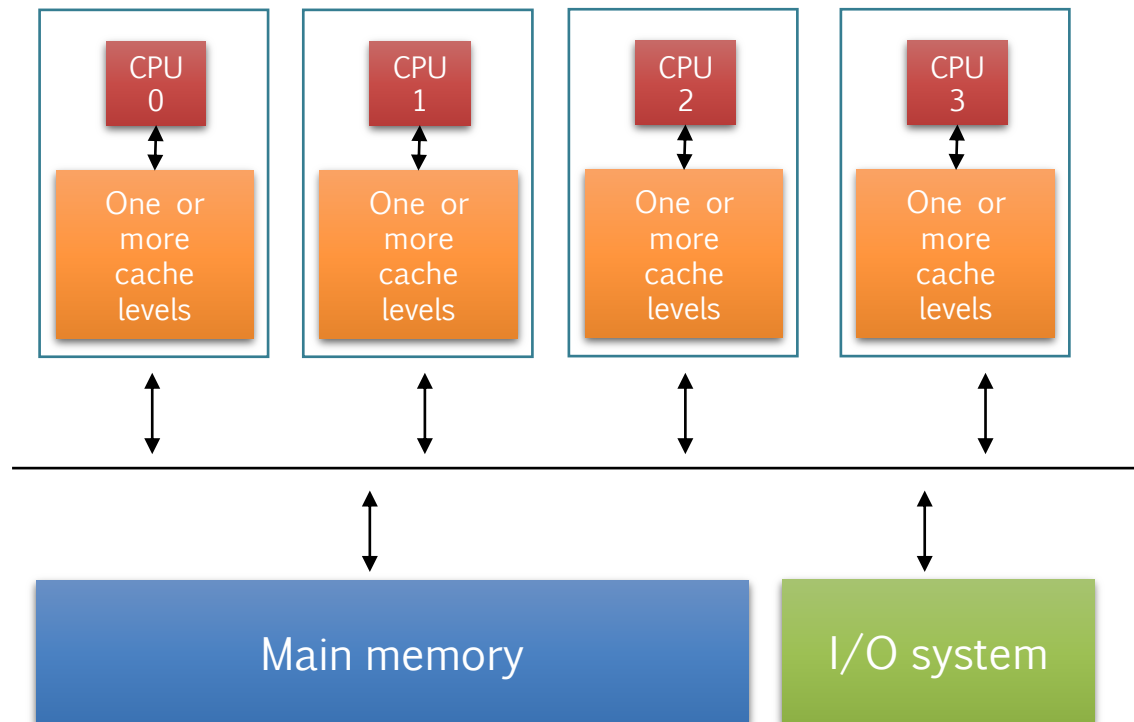




# Symmetric multi-processing

- › Memory: centralized with bus interconnect, I/O
- › Typically, **multi**-core (sub)systems
  - Examples: Sun Enterprise 6000, SGI Challenge, Intel (this laptop)

Can be 1 bus, N  
busses, or any  
network

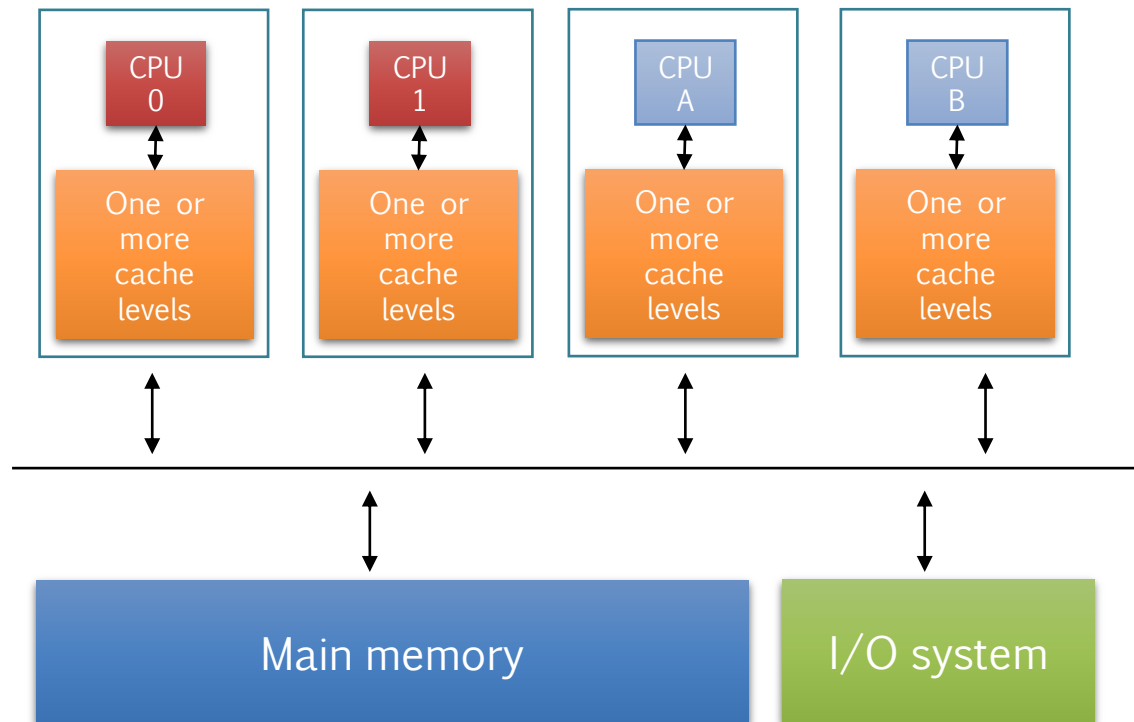




# Asymmetric multi-processing

- › Memory: centralized with uniform access time (UMA) and bus interconnect, I/O
- › Typically, multi-core (sub)systems
  - Examples: ARM Big.LITTLE, NVIDIA Tegra X2 (Drive PX)

Can be 1 bus, N  
busses, or any  
network



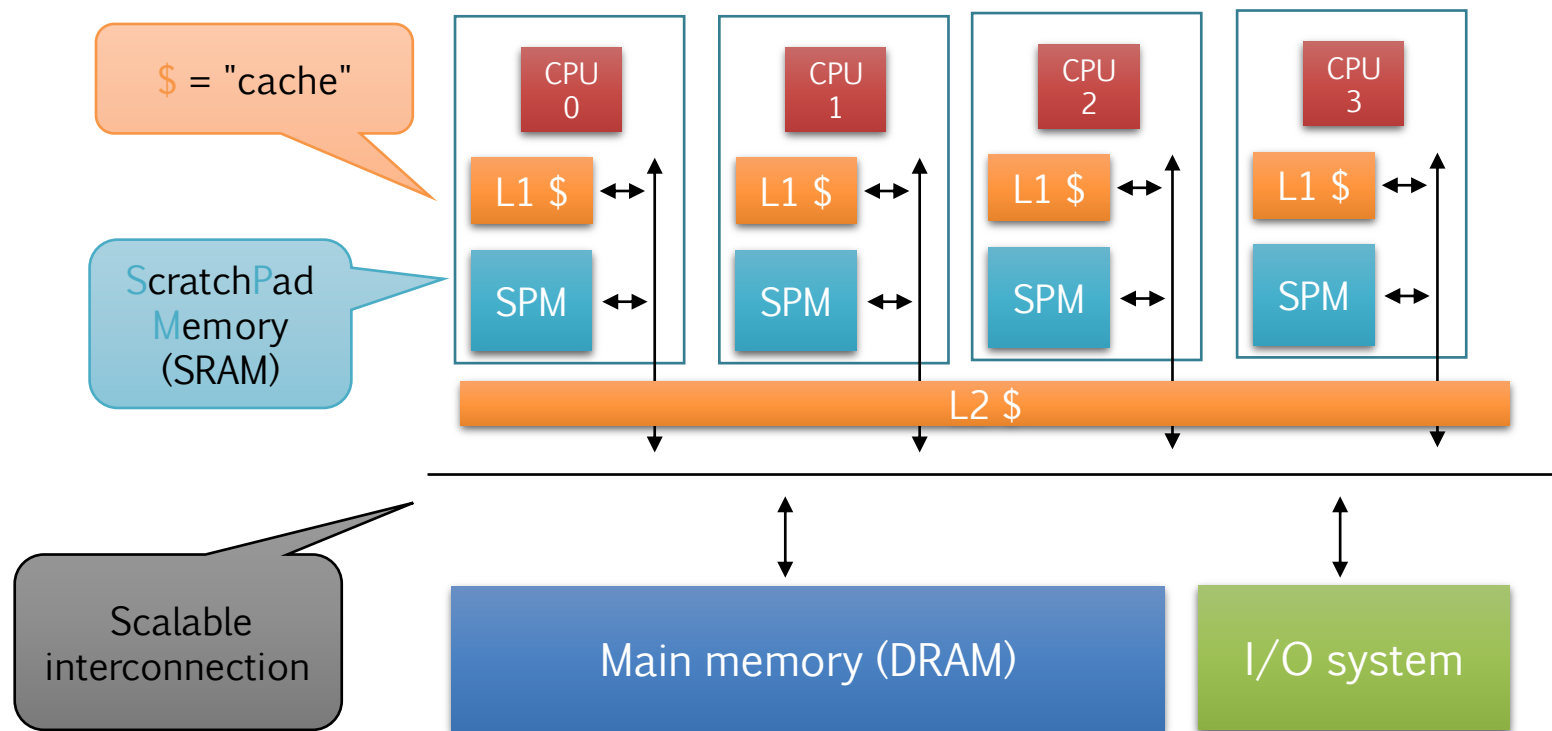


# SMP – distributed shared memory

› Non-Uniform Access Time - **NUMA**

› **Scalable** interconnect

- Typically, **many** cores
- Examples: embedded accelerators, GPUs

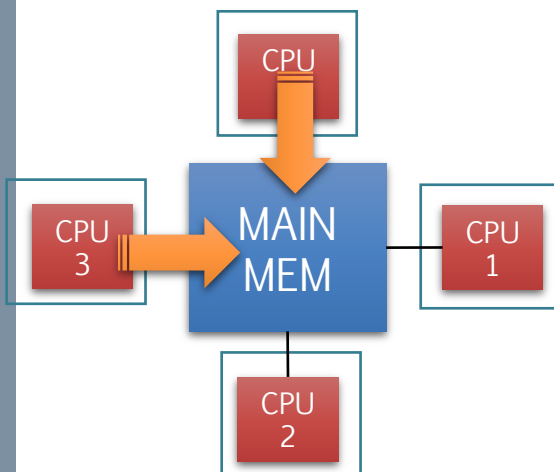




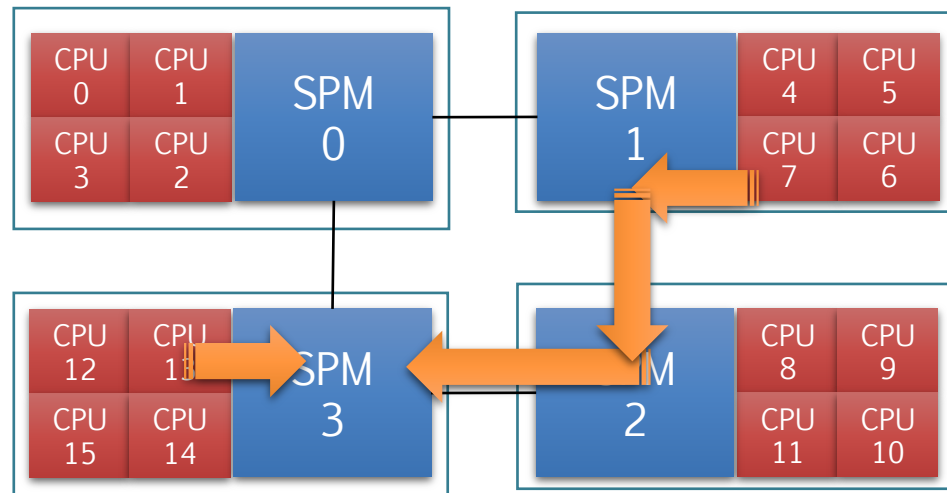
# UMA vs. NUMA

- › Shared mem: every thread can access every memory item
  - (Not considering security issues...)
- › Uniform Memory Access (UMA) vs Non-Uniform Memory Access (NUMA)
  - Different access time for accessing different memory spaces

## UMA



## NUMA







# UMA vs. NUMA

> Shared

– (

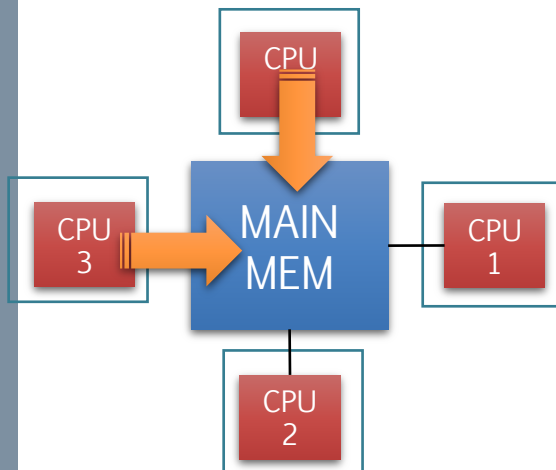
> Uniform

– [

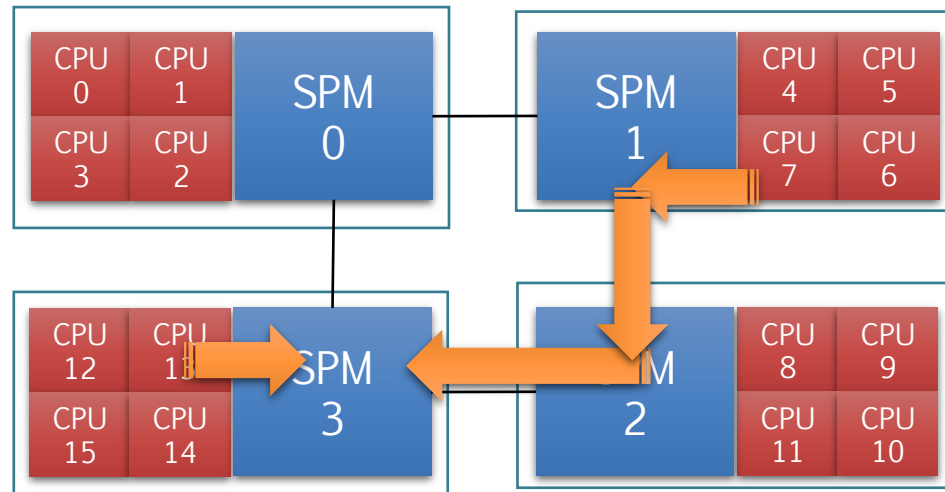
	MEM0	MEM1	MEM2	MEM3
CPU0...3	0 clock	10 clock	20 clock	10 clock
CPU4...7	10 clock	0 clock	10 clock	20 clock
CPU8...11	20 clock	10 clock	0 clock	10 clock
CPU12..15	10 clock	20 clock	10 clock	0 clock

(NUMA)

UMA



NUMA





# Programming abstractions



# (Simplest) threading model

---

## Fork-join execution model

- › The main, single thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                             &myattr,
                             my_pthread_fn,
                             NULL);
    // Here, the main thread can do other stuff!

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```



# (Simplest) threading model

---

## Fork-join execution model

- › The main, single thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                             &myattr,
                             my_pthread_fn,
                             NULL);
    // Here, the main thread can do other stuff!

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





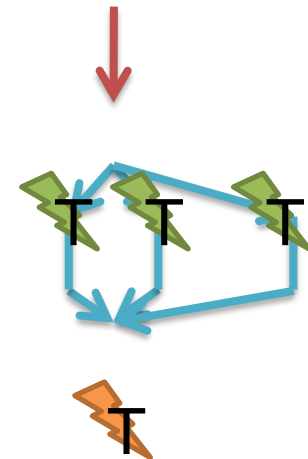
# (Simplest) threading model

## Fork-join execution model

- › The main, single thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                             &myattr,
                             my_pthread_fn,
                             NULL);
    // Here, the main thread can do other stuff!

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





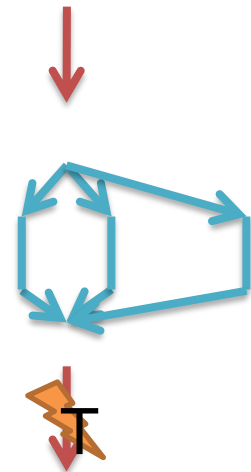
# (Simplest) threading model

## Fork-join execution model

- › The main, single thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                             &myattr,
                             my_pthread_fn,
                             NULL);
    // Here, the main thread can do other stuff!

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





# (Simplest) threading model

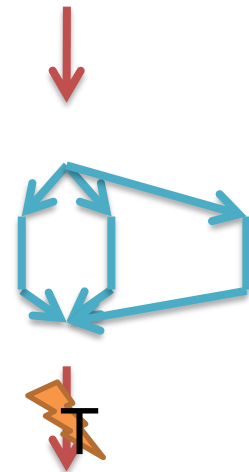
## Fork-join execution model

- › The main, single thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // == 0
                             &myattr,
                             my_pthread,
                             NULL);
    // Here, the main thread can do o

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &return_val); // <== JOIN
}
```

Let's see  
this in  
action





# Master-slave threading model

---

## Fork-join execution model

- › The main, **Master** thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                             &myattr,
                             my_pthread_fn,
                             NULL);

    // Here, the main thread can do other stuff!
    other_fn();

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





# Master-slave threading model

## Fork-join execution model

- › The main, **Master** thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                               &myattr,
                               my_pthread_fn,
                               NULL);

    // Here, the main thread can do other stuff!
    other_fn();

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





# Master-slave threading model

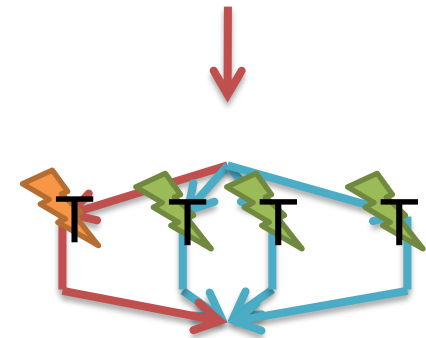
## Fork-join execution model

- › The main, **Master** thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                               &myattr,
                               my_pthread_fn,
                               NULL);

    // Here, the main thread can do other stuff!
    other_fn();

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





# Master-slave threading model

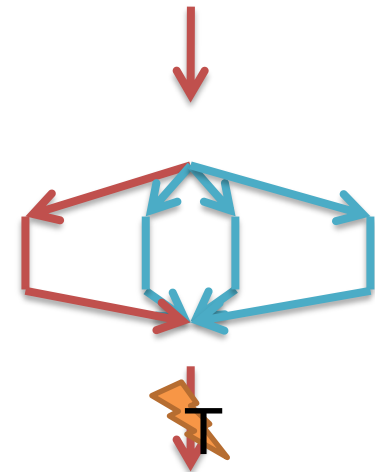
## Fork-join execution model

- › The main, **Master** thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                              &myattr,
                              my_pthread_fn,
                              NULL);

    // Here, the main thread can do other stuff!
    other_fn();

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```





# Work partitioning

---

Several models, here to cite a few

- › Data parallelism (see also GPGPUS)
  - We're getting there, don't worry...
- › Reduction
- › Task parallelism (aka: work queue)
- › ...



# Data parallelism

---

(Aka: data decomposition, loop decomposition, SPMD, SIMD\*...)

Parallel threads execute the same operation(s) on multiple data

- › Data is typically an array, a matrix (image)....
  - Note: you typically map the iteration id of the loop to the data index
- › **Partitioning strategy** defines how many iterations (**chunk**) every thread will perform
  - From 1 iteration, to loop size





# Exercise

Let's  
code!

Create an array of  $N$  elements

- › Put inside each array element its index, multiplied by '2'
- › `arr[0] = 0; arr[1] = 2; arr[2] = 4; ...and so on..`

Now, do it in parallel with a team of  $T$  PThreads

- › Assume  $N$  is a multiple of  $T$
- › “Decompose” the `for` construct, so that every thread manages (chunk size is)  $N/T$  iterations



# Reduction

wikipedia

*The reduction clause can be used to perform some forms of **recurrence** calculations (involving **mathematically associative and commutative operators**) in parallel. For parallel [...], a **private** copy of each list item is created, one for each implicit task, as if the private clause had been used. [...] The private copy is then initialized as specified above. At the end of the region for which the reduction clause was specified, the original list item is updated by **combining its original value with the final value of each of the private copies**, using the combiner of the specified reduction-identifier.*

E.g., average value of a sequence (array/vector)

- › Create a thread-local copy of a variable
- › Accumulate sums only for the assigned part of the array/vector
- › Then, sum the partial sums (and divide by size)



# Exercise

---

Let's  
code!

Create an array of  $N$  elements

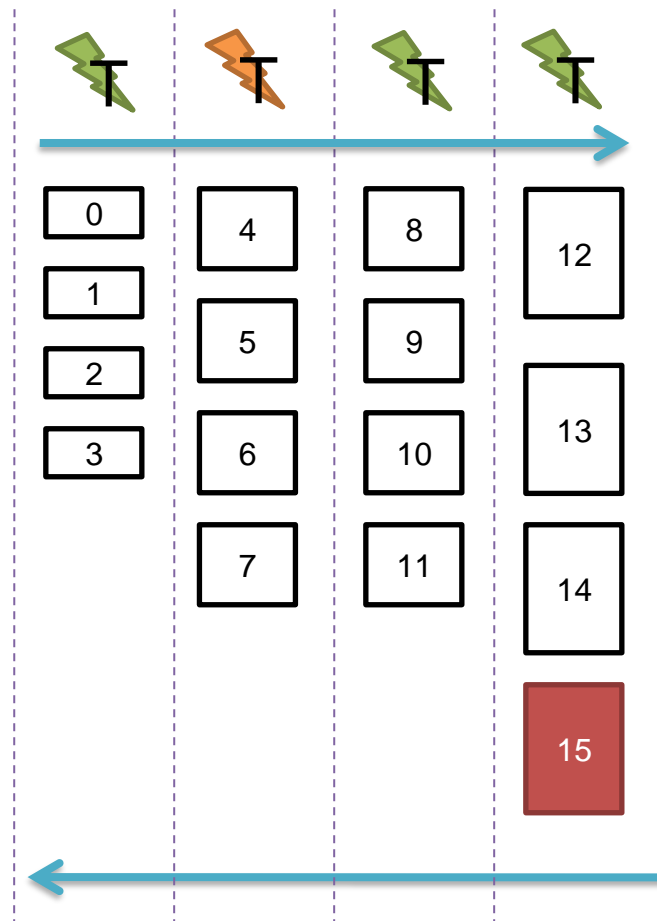
- › ..or a vector
- › Initiate it randomly
- › Now, compute its average value using multi-treading and reduction paradigm





# Unbalanced loop partitioning

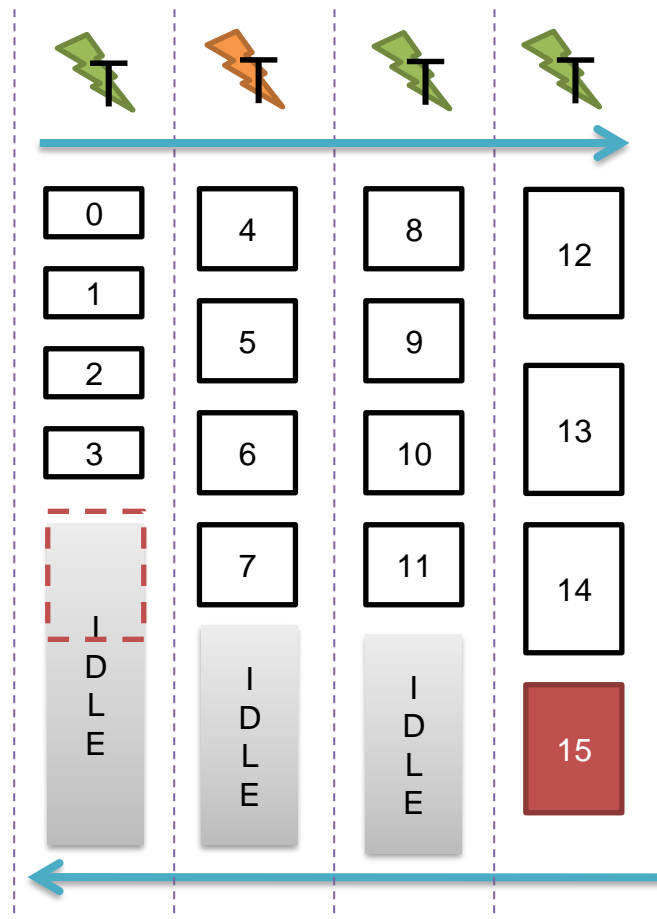
- › So far, we assigned iterations «statically»
  - Might not be effective nor efficient





# Unbalanced loop partitioning

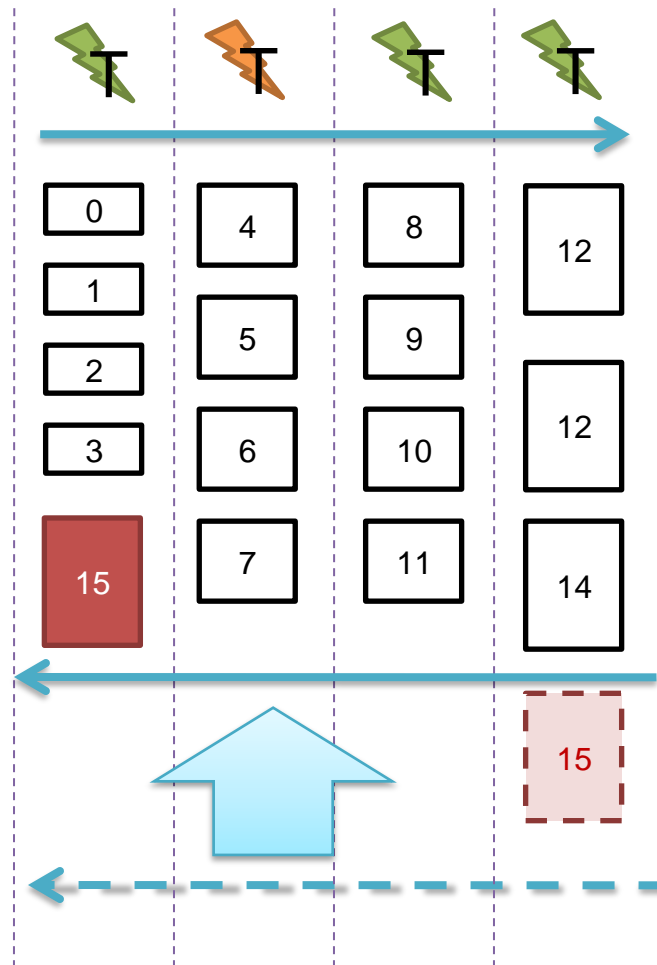
- › So far, we assigned iterations «statically»
  - Might not be effective nor efficient





# How can we manage dynamics/irregular workloads?

› We would like something like this..



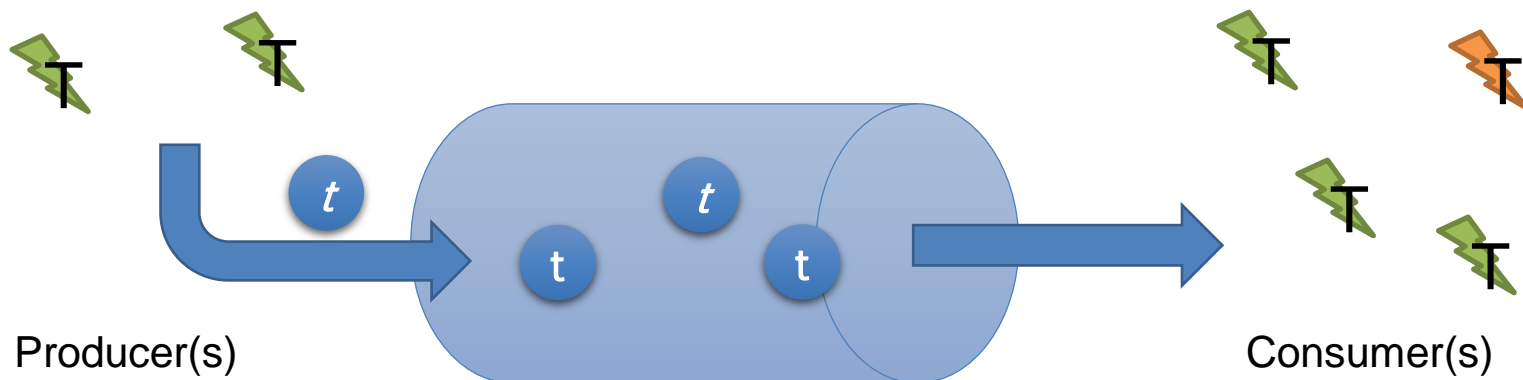


# A different parallel paradigm

Implements a **producer-consumer paradigm**

Managed by a **task queue**

- › Where units of work (**tasks**)
- › are pushed by threads (**q\_push** primitive)
- › and pulled and executed by threads (**q\_pop** primitive)



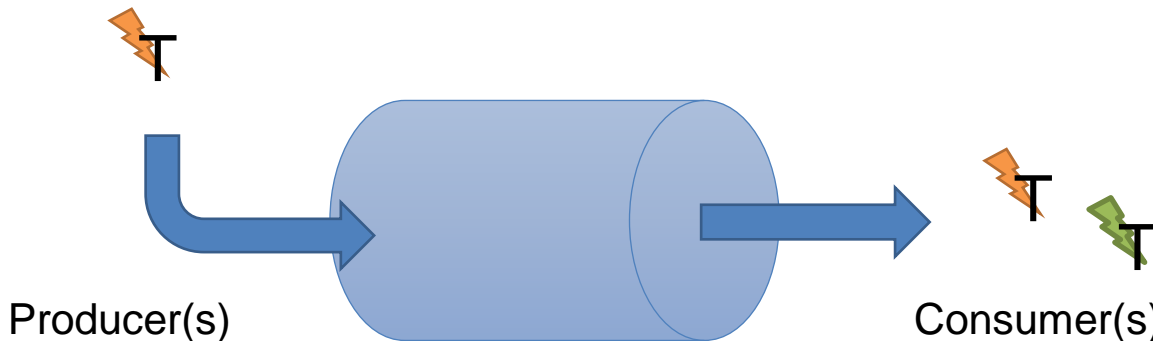


# «What» happens «when»?

```
void t0() {  
    // Task 0  
}  
void t1() {  
    // Task 1 pushes t2 in the q  
    q_push(t2());  
}  
void t2() {  
    // Task 2  
}  
  
void thread_fn() {  
    // Push t0 and t1  
    q_push(t0());  
    q_push(t1());  
}
```

```
void other_thread_fn() {  
    // Pop a task (which one?)  
    t = q_pop();  
    // Execute it  
    t();  
}
```

\*pseudo-code



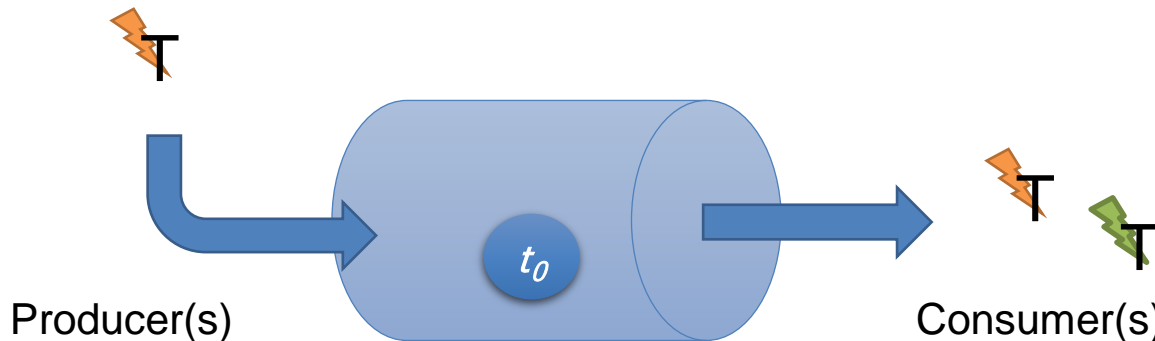


# «What» happens «when»?

```
void t0() {  
    // Task 0  
}  
void t1() {  
    // Task 1 pushes t2 in the q  
    q_push(t2());  
}  
void t2() {  
    // Task 2  
}  
  
void thread_fn() {  
    // Push t0 and t1  
    q_push(t0());  
    q_push(t1());  
}
```

```
void other_thread_fn() {  
    // Pop a task (which one?)  
    t = q_pop();  
    // Execute it  
    t();  
}
```

\*pseudo-code



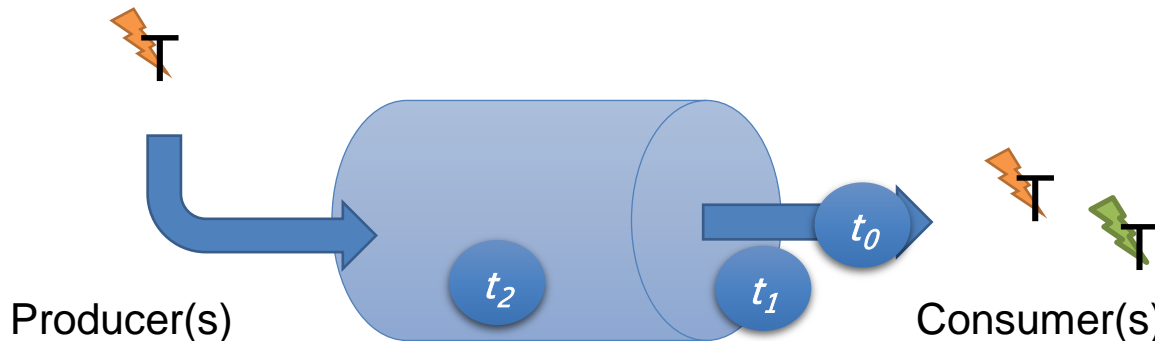


# «What» happens «when»?

```
void t0() {  
    // Task 0  
}  
void t1() {  
    // Task 1 pushes t2 in the q  
    q_push(t2());  
}  
void t2() {  
    // Task 2  
}  
  
void thread_fn() {  
    // Push t0 and t1  
    q_push(t0());  
    q_push(t1());  
}
```

```
void other_thread_fn() {  
    // Pop a task (which one?)  
    t = q_pop();  
    // Execute it  
    t();  
}
```

\*pseudo-code

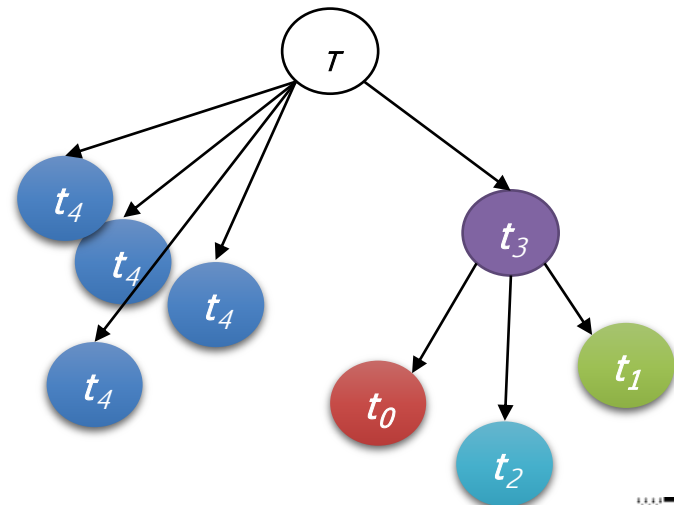
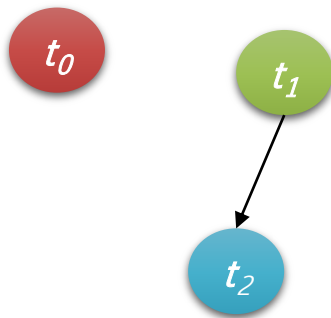




# The queue

It's a shared resource

- › Its primitives `q_push` and `q_pop` are *thread-safe*, i.e., their concurrent access is protected by semaphores
- › Typically implemented as a FIFO queue
- › ..but we can also have more complex semantics (e.g., parent-son => DAGs) among tasks





# References

---



## Course website

- › [http://hipert.unimore.it/people/paolob/pub/Industrial\\_Informatics/index.html](http://hipert.unimore.it/people/paolob/pub/Industrial_Informatics/index.html)

## My contacts

- › [paolo.burgio@unimore.it](mailto:paolo.burgio@unimore.it)
- › <http://hipert.mat.unimore.it/people/paolob/>

## Resources

- › "Parallel programming" course by "a guy" @UNIMORE
  - [https://hipert.unimore.it/people/paolob/pub/Calcolo\\_Parallelo/index.html](https://hipert.unimore.it/people/paolob/pub/Calcolo_Parallelo/index.html)
  - <https://github.com/HiPeRT/cp19/>
- › A "small blog"
  - <http://www.google.com>